

HOW TO SCALE RELATIONAL DATABASE MIGRATIONS

Jumana Bahrainwala

@JumzB

Table "public.users"

Column	Type
id	integer
active	boolean
first	string
last	string
email	string

TIME



TIME

```
BEGIN;  
  
ALTER TABLE  
    USERS  
    RENAME  
    COLUMN  
        first TO  
        first_name;
```

Table "public.users"

Column	Type
id	integer
active	boolean
first	string
last	string
email	string

TIME

```
BEGIN;  
  
ALTER TABLE  
    USERS  
    RENAME  
    COLUMN  
        first TO  
        first_name;
```

```
BEGIN;  
  
Select * from  
    Users where  
        id=2;  
  
COMMIT;
```

Table "public.users"

Column	Type
id	integer
active	boolean
first	string
Last	string
email	string

TIME

```
BEGIN;  
  
ALTER TABLE  
    USERS  
    RENAME  
    COLUMN  
        first TO  
        first_name;
```

```
BEGIN;  
  
Select * from  
    Users where  
        id=2;  
  
COMMIT;
```

Table "public.users"

Column	Type
id	integer
active	boolean
first	string
Last	string
email	string

PID	USER	PRI	NI	VIRT	RES	S	CPU%	MEM%	TIME+	Command
4511	root	32	0	0	0	?	0.0	0.0	0:00.00	login
4482	jzbahrai	24	0	4373M	3612	?	0.0	0.0	0:00.00	postgres: jzbahrai jzbahrai ::1(56893) SELECT waiting
4481	jzbahrai	25	0	4190M	5444	?	0.0	0.0	0:00.01	psql -h localhost -U jzbahrai

TIME

```
BEGIN;  
  
ALTER TABLE  
    USERS  
    RENAME  
    COLUMN  
        first TO  
        first_name;
```

```
COMMIT;
```

```
BEGIN;  
  
Select * from  
    Users where  
        id=2;  
  
COMMIT;
```

Table "public.users"

Column	Type
id	integer
active	boolean
first	string
Last	string
email	string

PID	USER	PRI	NI	VIRT	RES	S	CPU%	MEM%	TIME+	Command
4511	root	32	0	0	0	?	0.0	0.0	0:00.00	login
4482	jzbahrai	24	0	4373M	3612	?	0.0	0.0	0:00.00	postgres: jzbahrai jzbahrai ::1(56893) SELECT waiting
4481	jzbahrai	25	0	4190M	5444	?	0.0	0.0	0:00.01	psql -h localhost -U jzbahrai

TIME

```
BEGIN;  
  
ALTER TABLE  
    USERS  
    RENAME  
    COLUMN  
        first TO  
        first_name;
```

```
COMMIT;
```

```
BEGIN;  
  
Select * from  
    Users where  
        id=2;  
  
COMMIT;
```

Table "public.users"

Column	Type
id	integer
active	boolean
first	string
Last	string
email	string

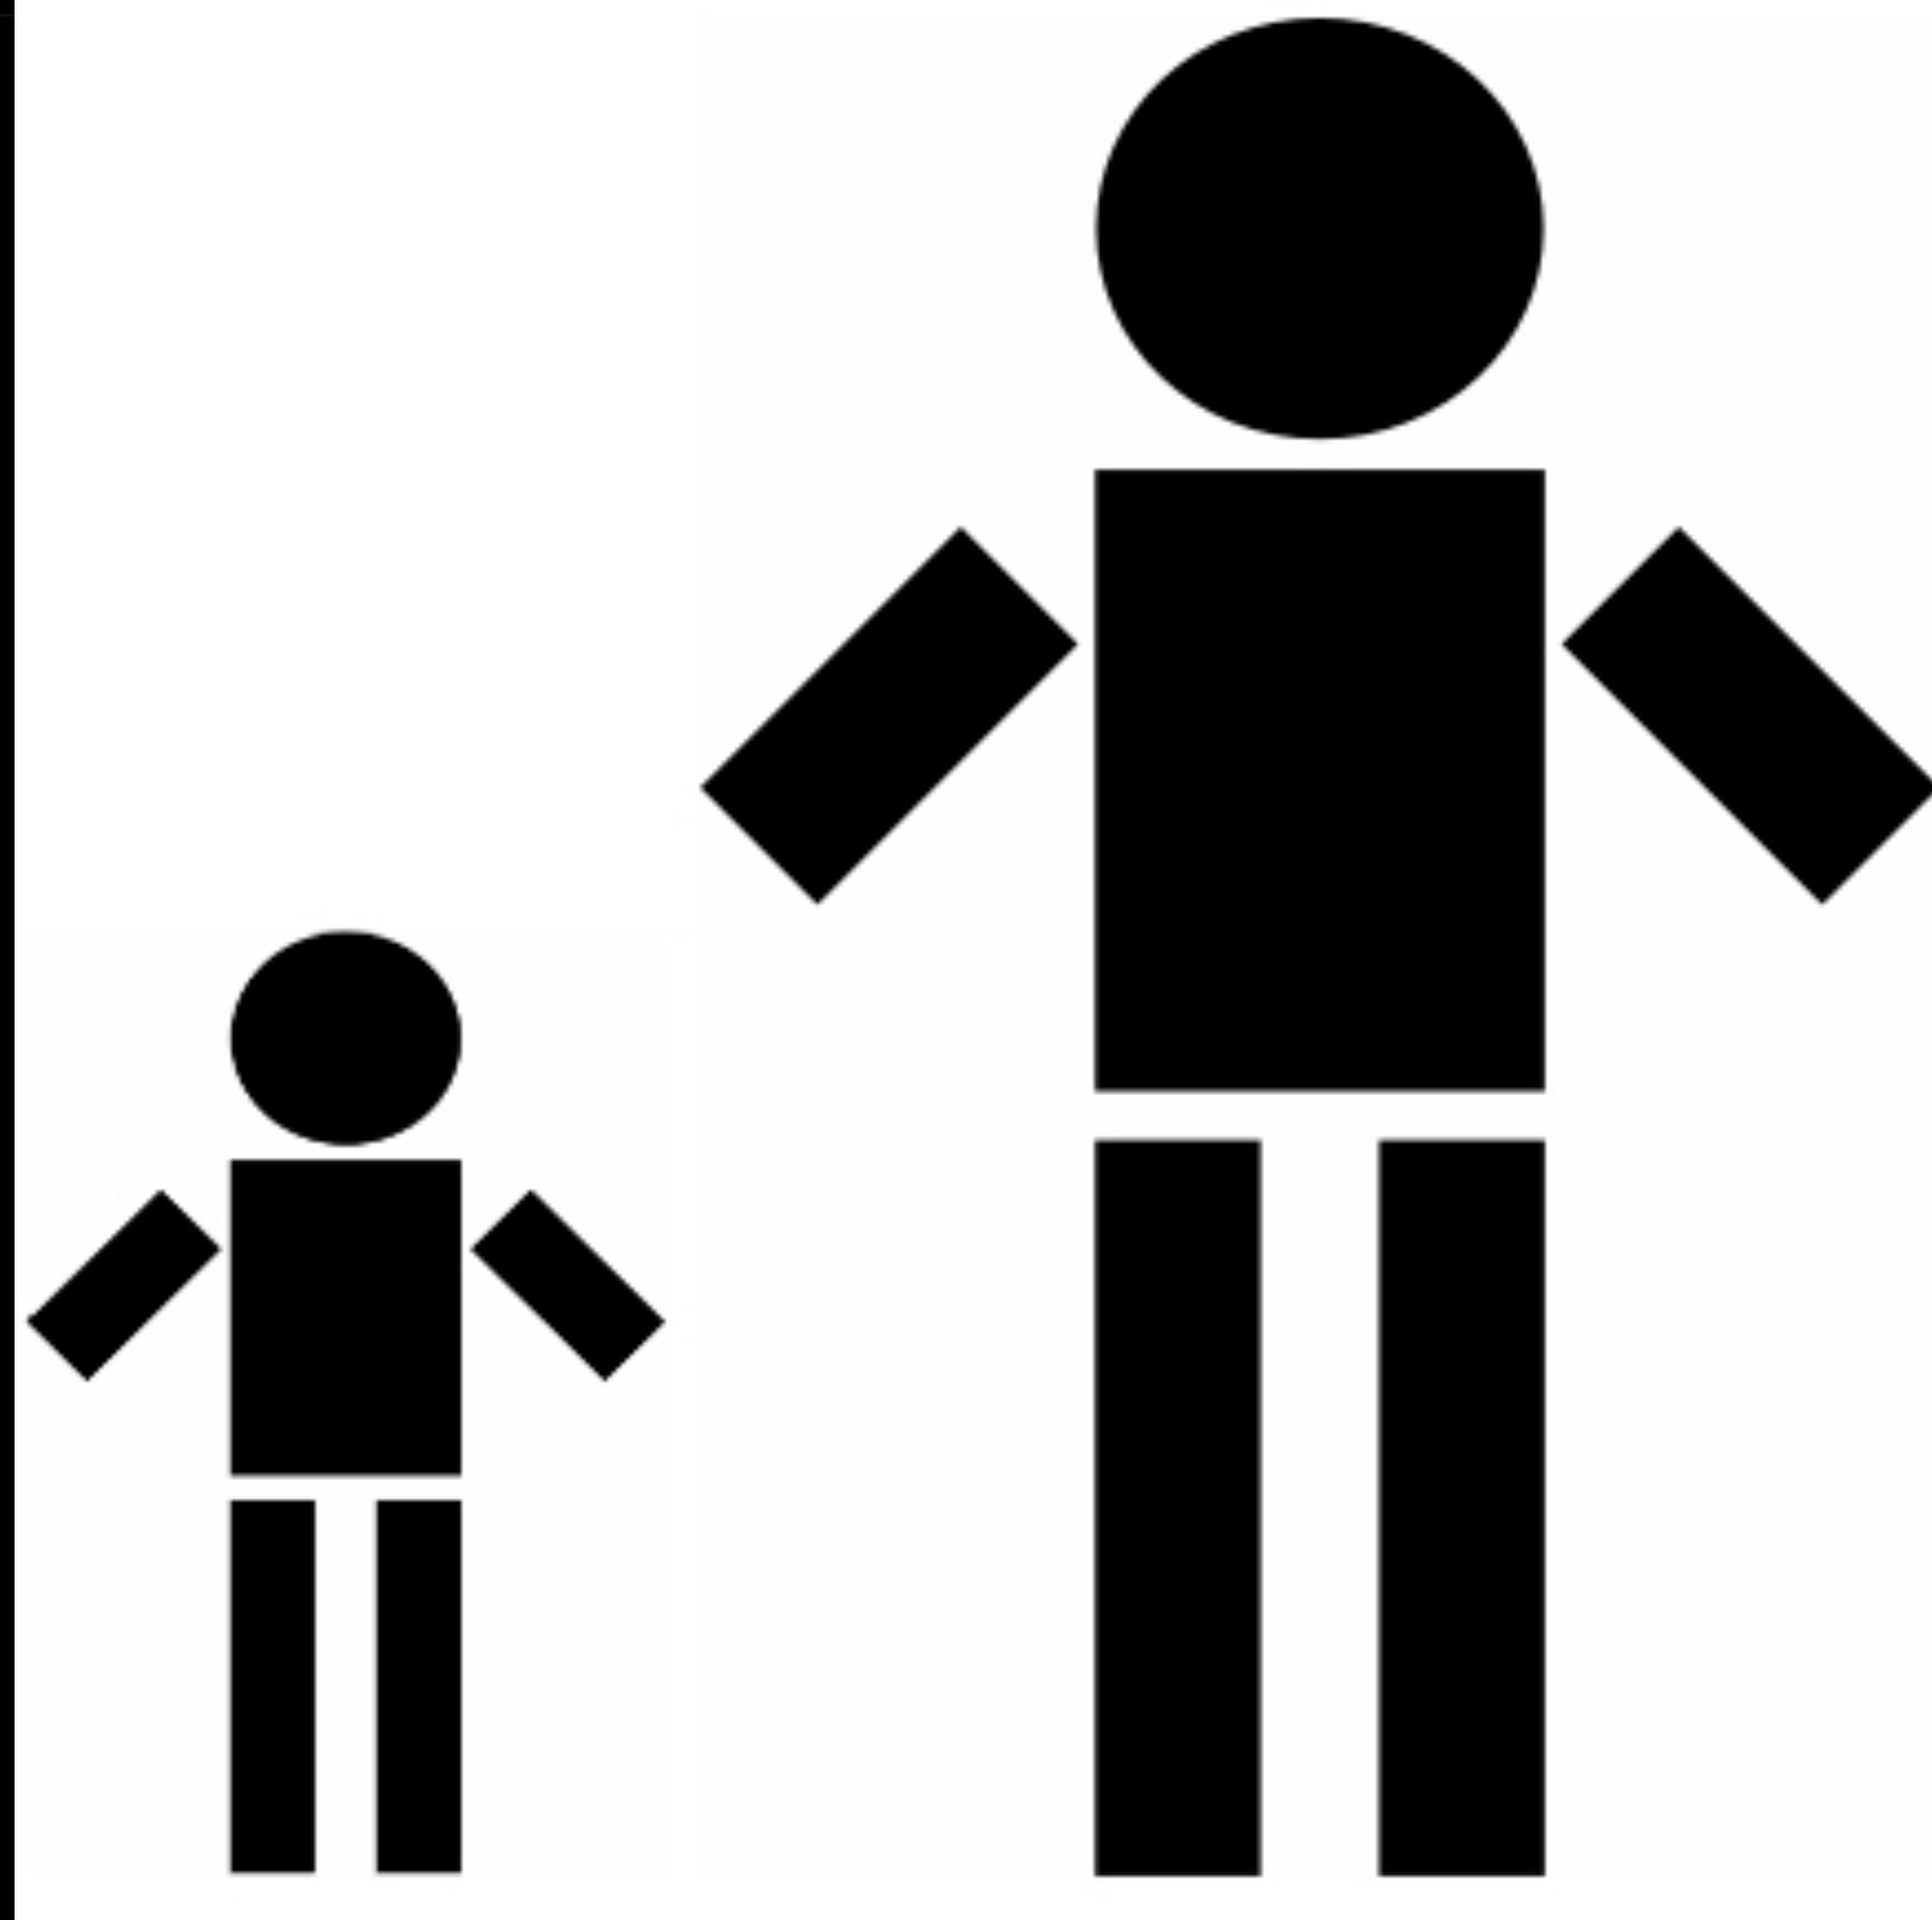
Wait...what?

PID	USER	PRI	NI	VIRT	RES	S	CPU%	MEM%	TIME+	Command
4511	root	32	0	0	0	?	0.0	0.0	0:00.00	login
4482	jzbahrai	24	0	4373M	3612	?	0.0	0.0	0:00.00	postgres: jzbahrai jzbahrai ::1(56893) SELECT waiting
4481	jzbahrai	25	0	4190M	5444	?	0.0	0.0	0:00.01	psql -h localhost -U jzbahrai

WHEN YOU CHANGE A COLUMN TYPE

AND LOCK YOUR DATABASE FOR 10
MINUTES





“Most of these examples are Postgres specific,
but the concepts apply to most SQL like DBs.”

scaling with tools

scaling with people

remaining issues getting to zero
downtime

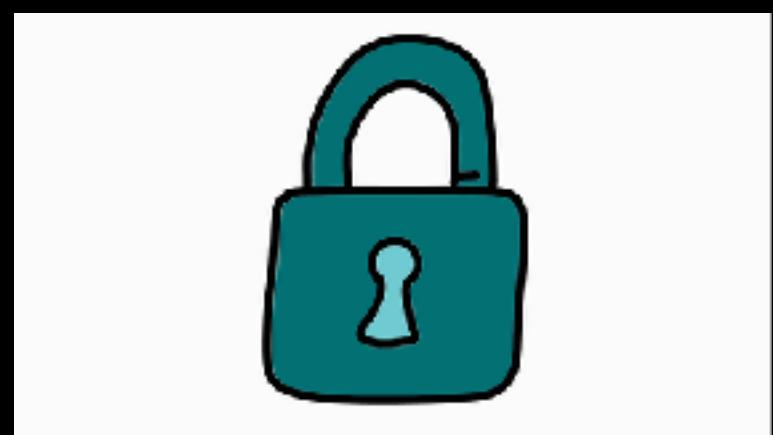
Don't lock yourself
out

scaling with tools

don't lock yourself out



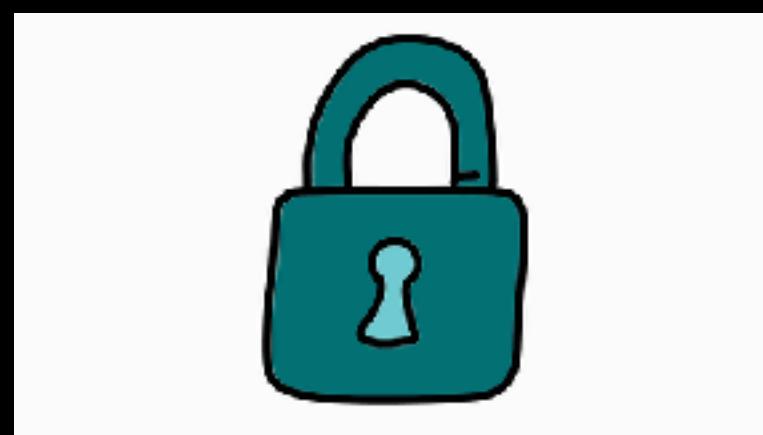
don't lock yourself out



don't lock yourself out



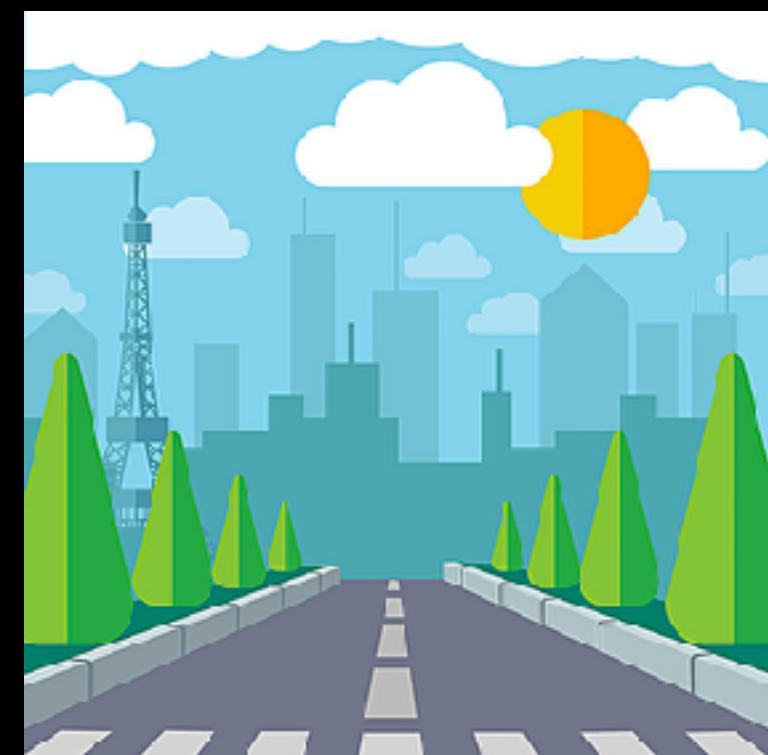
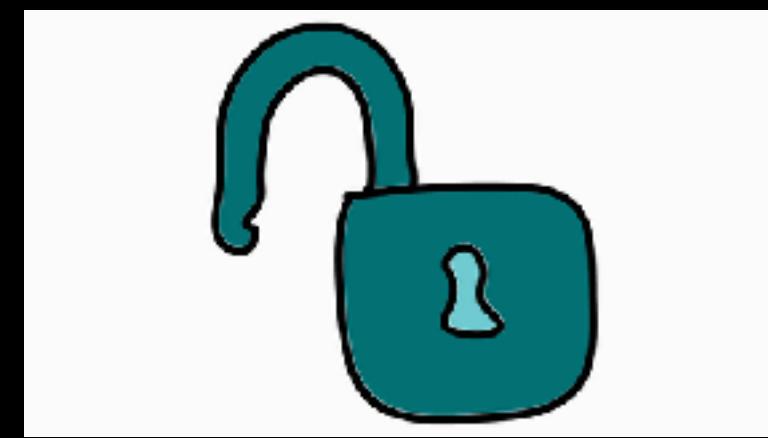
don't lock yourself out



don't lock yourself out



don't lock yourself out



don't lock yourself out



ACCESS EXCLUSIVE

Alter Table
Drop Table
Truncate
Reindex
Cluster
Vacuum

don't lock yourself out



ROW SHARE

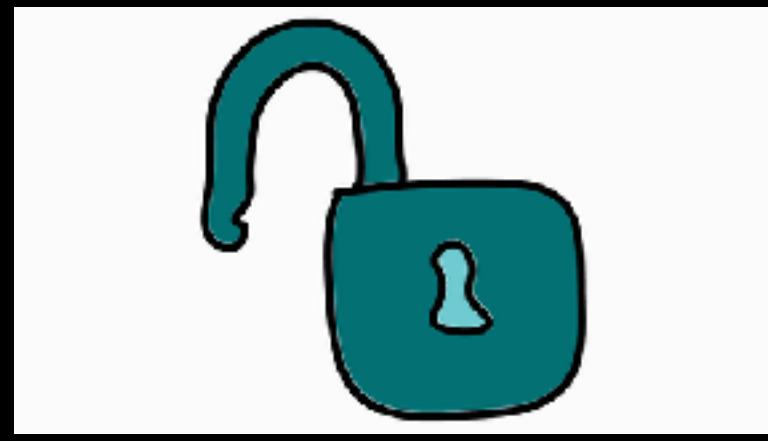
Select for Update
Select for Share



ROW EXCLUSIVE

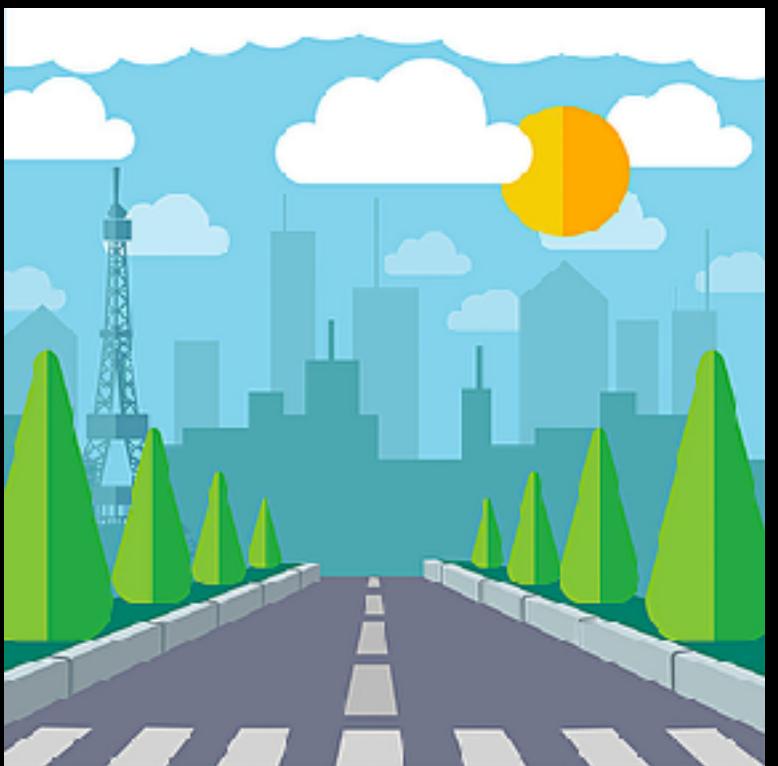
Update, Delete, Insert
Anything that modifies Data.

don't lock yourself out



ACCESS SHARE

Select



don't lock yourself out

READS

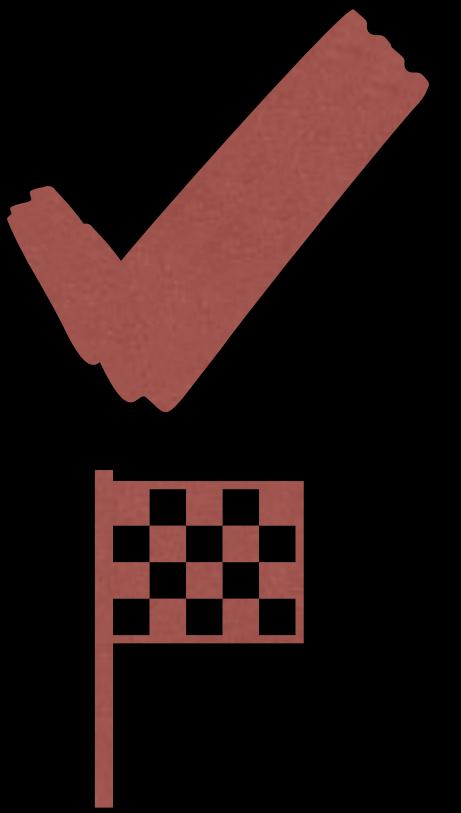


READS

don't lock yourself out

READS

READS



READS

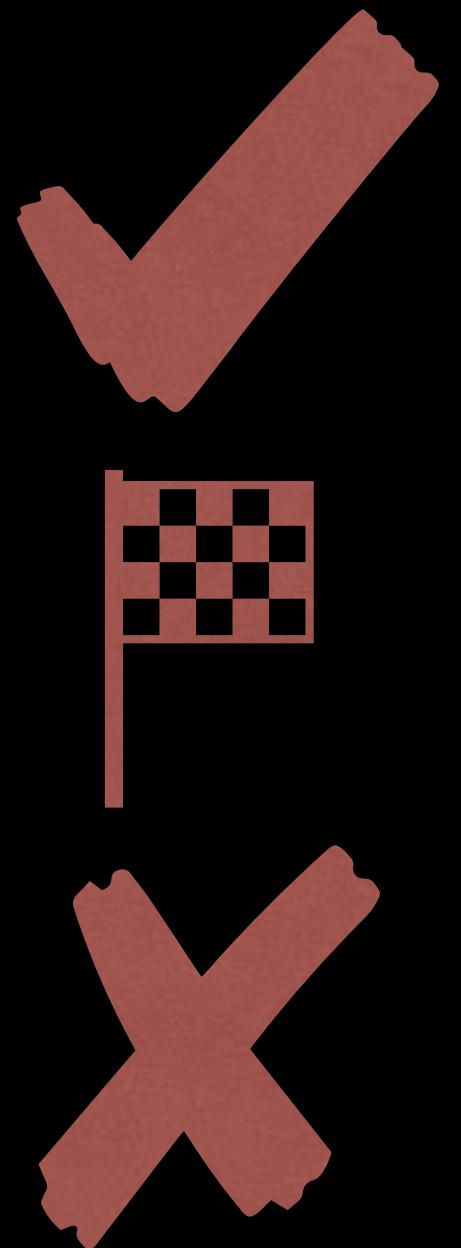
WRITES

don't lock yourself out

READS

READS

READS,
WRITES



READS

WRITES

DDL CHANGES

“Look for locks that are conflicting.

Conflicting locks are what cause downtime”

Don't lock yourself
out

scaling with tools

scaling with tools

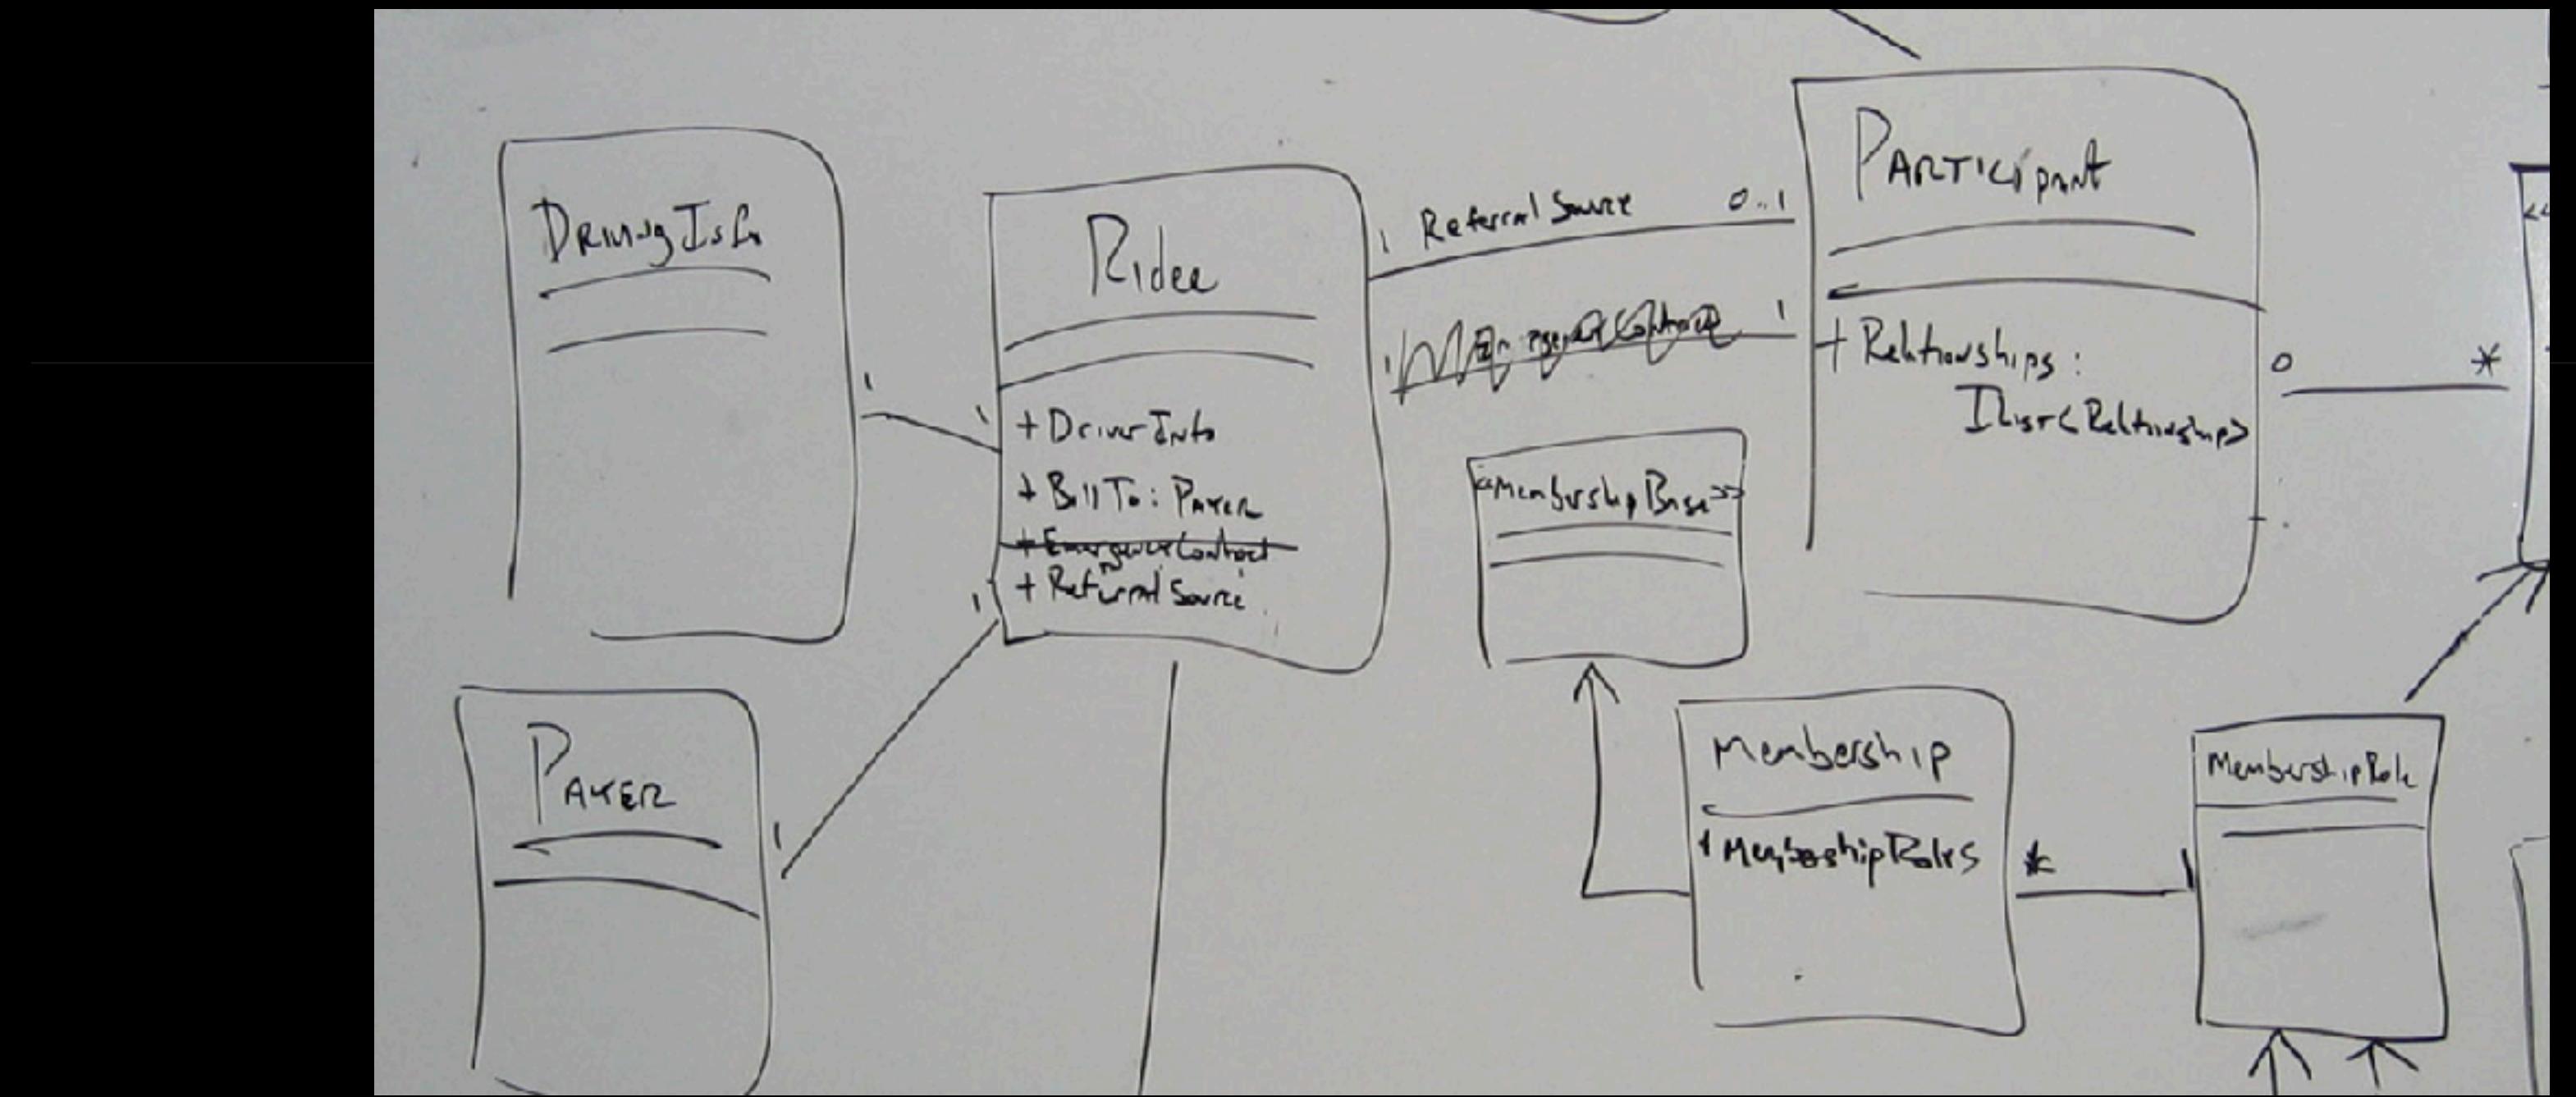
Don't lock yourself
out

Break your
migrations

Schema Changes

Schema Changes

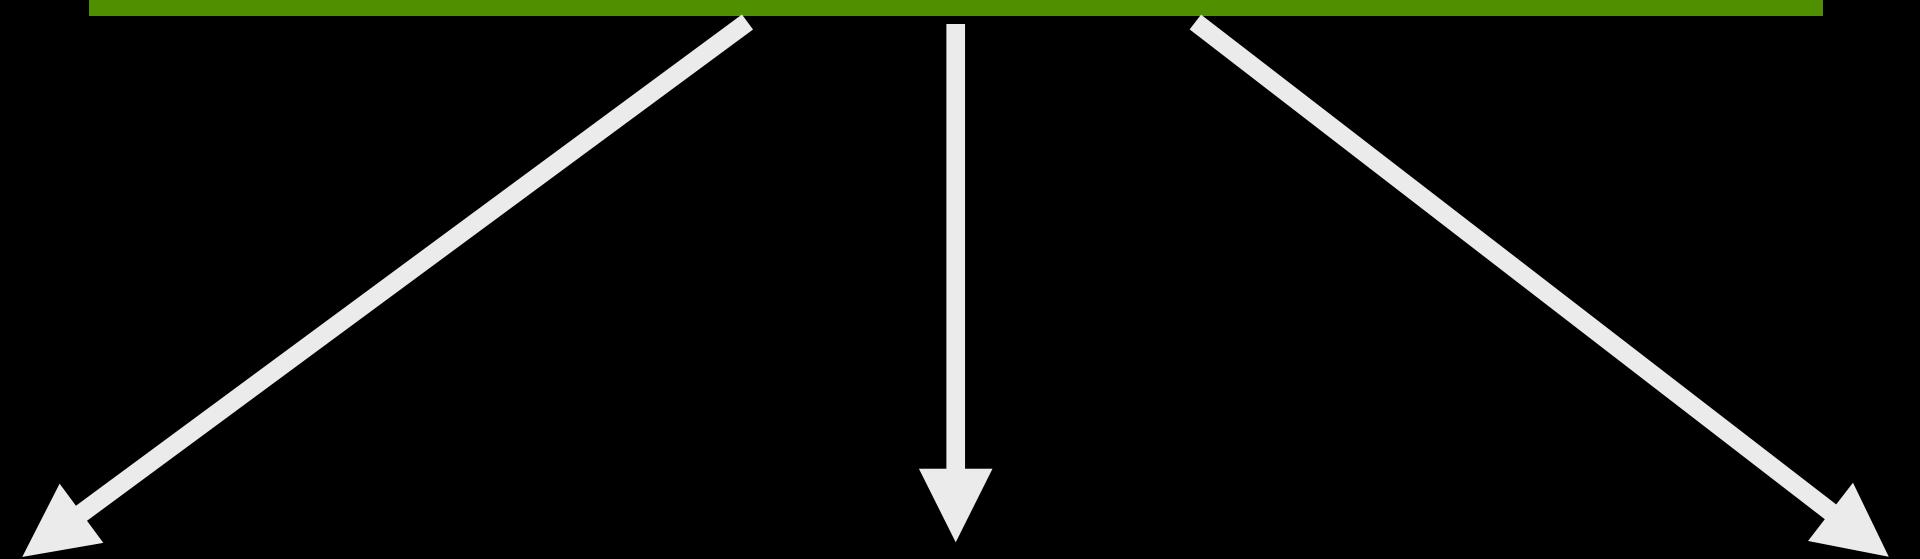




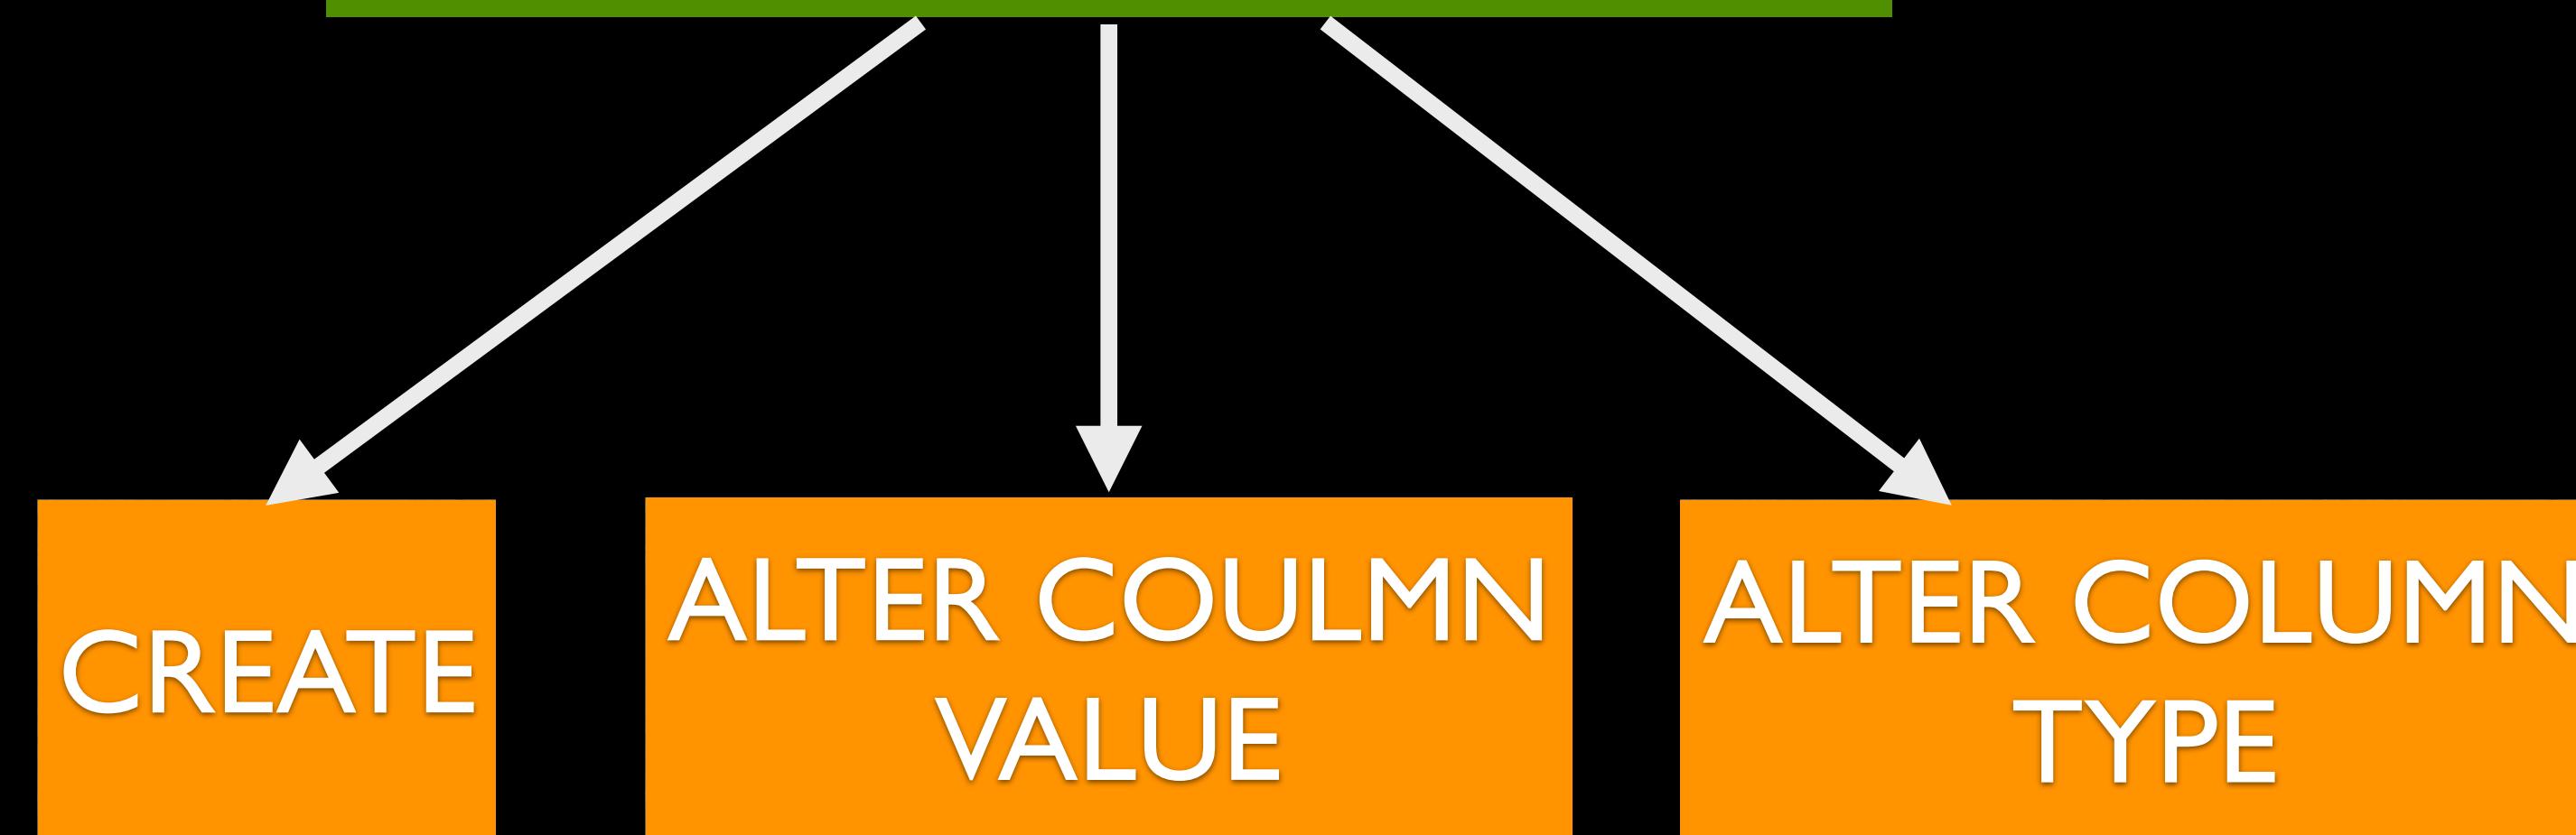
Data Definition Language

DATA DEFINITION LANGUAGE

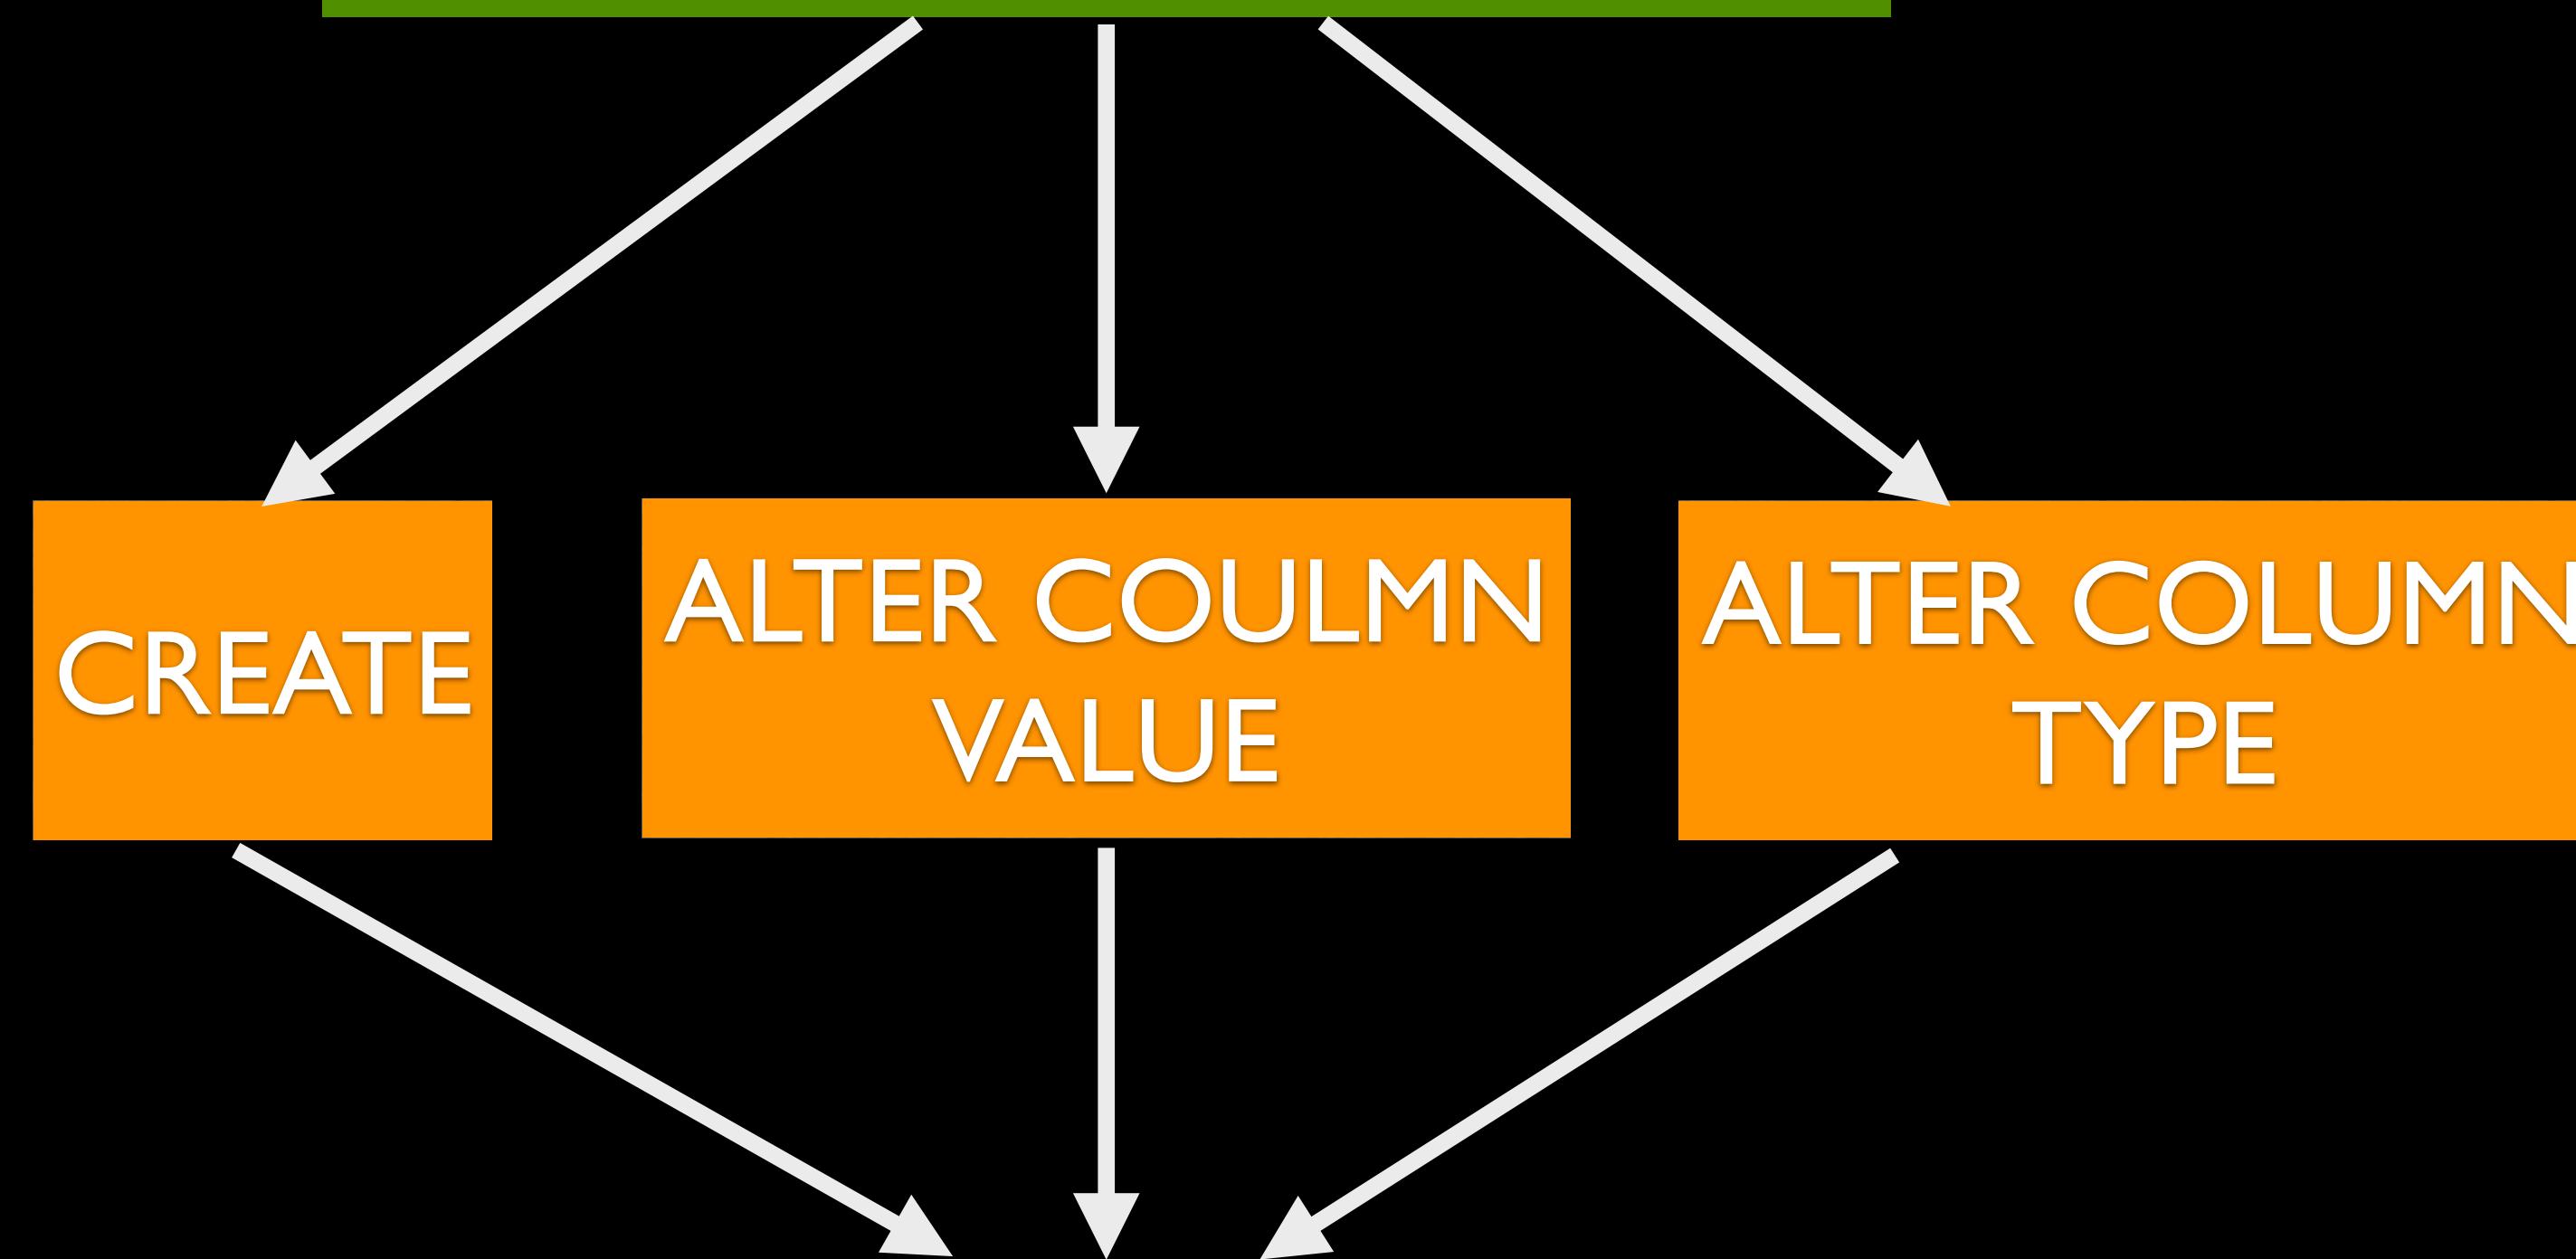
DATA DEFINITION LANGUAGE



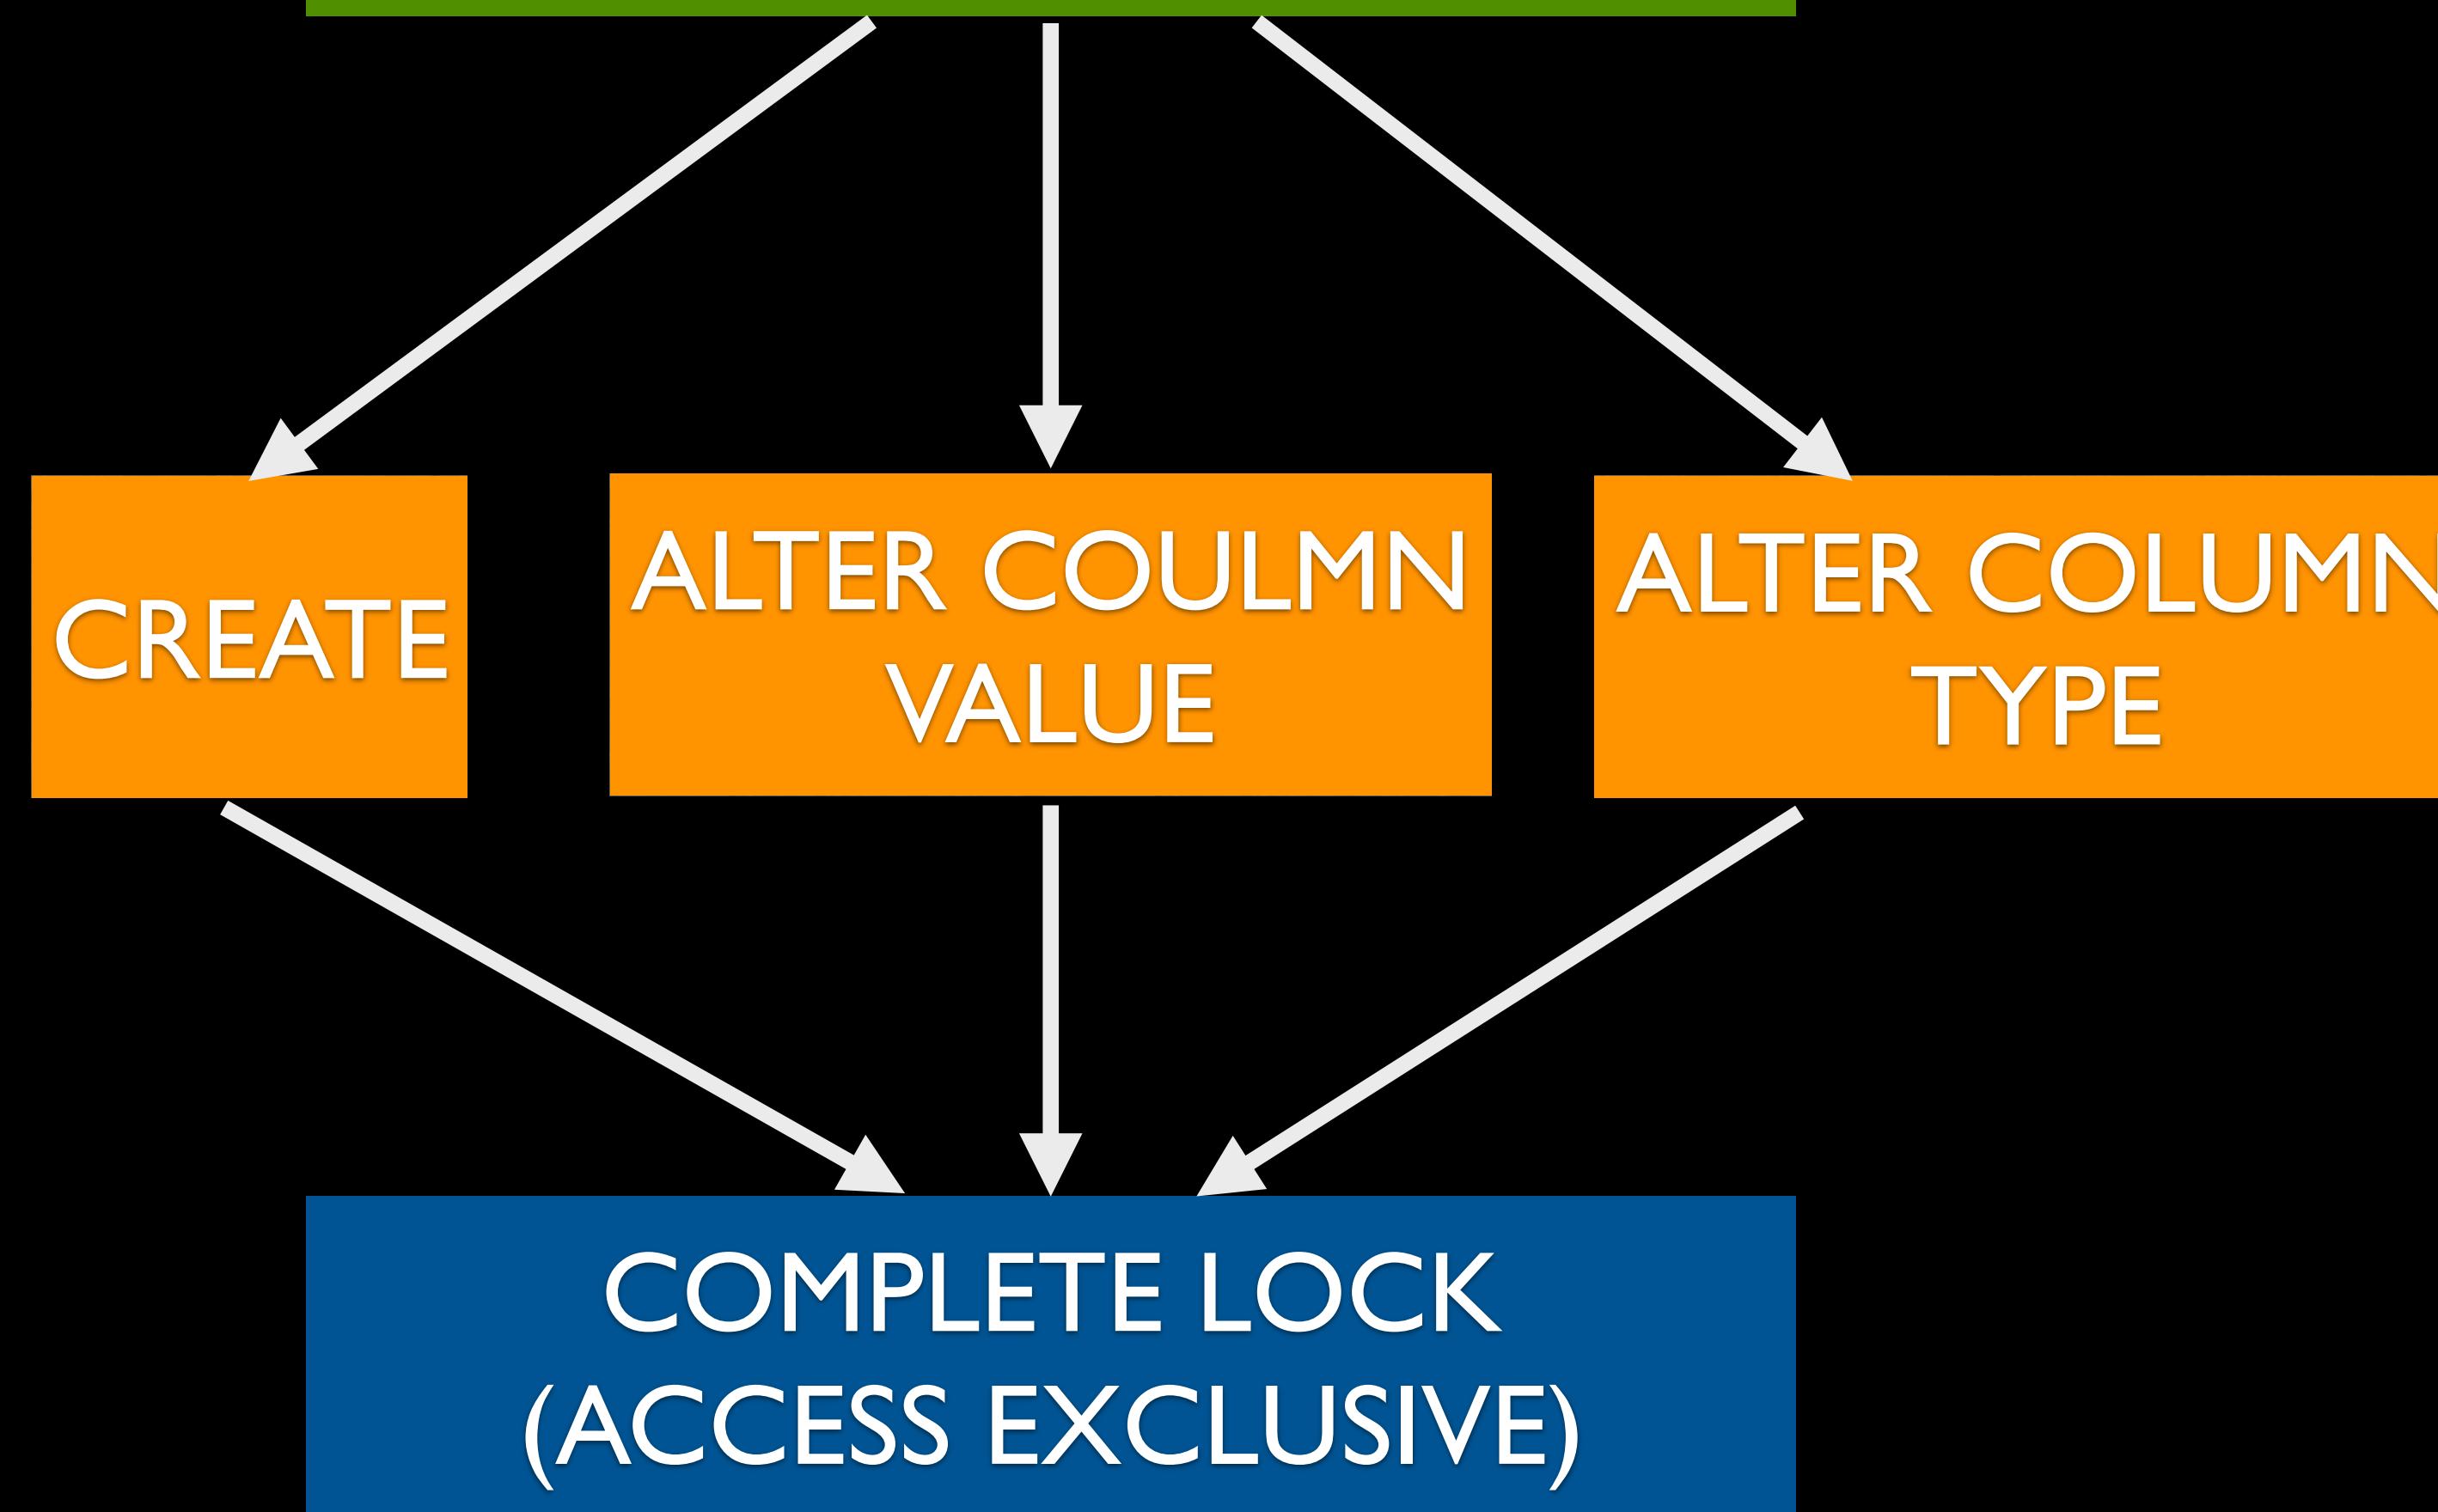
DATA DEFINITION LANGUAGE



DATA DEFINITION LANGUAGE



DATA DEFINITION LANGUAGE



break your migrations

Schema Changes

- QUICK
- SMALL TRANSACTIONS

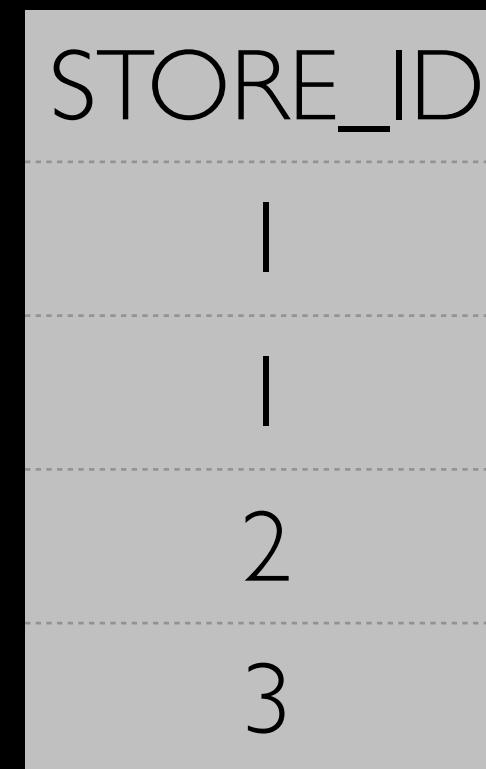


break your migrations

Index Changes

break your migrations

An **index** is a way to access some data quickly.



STORE_ID_INDEX

break your migrations

Index Changes

- CONCURRENTLY
- NON-CONCURRENT will LOCK (expensive)
- SHARE LOCK

break your migrations

Data Changes

break your migrations

Data Changes

Copy Data

Insert, Update
statements should
all be in their own
transactions

ROW EXCLUSIVE

break your migrations

Data Changes

Copy Data

Insert, Update
statements should
all be in their own
transactions

ROW EXCLUSIVE



Schema

```
CREATE TABLE
Users(
    id bigint
    first_name varchar(),
    last_name varchar(),
    email varchar(),
    phone varchar(),
    PRIMARY KEY (id)
);
```

Schema

```
CREATE TABLE  
Users(  
    id bigint  
    first_name varchar(),  
    last_name varchar(),  
    email varchar(),  
    phone varchar(),  
    PRIMARY KEY (id)  
);
```

Second Schema Migration

```
ALTER TABLE  
Users  
ADD COLUMN  
age_int int;  
  
ALTER TABLE  
Users  
ADD COLUMN  
status boolean
```

Schema

```
CREATE TABLE  
Users(  
    id bigint  
    first_name varchar(),  
    last_name varchar(),  
    email varchar(),  
    phone varchar(),  
    PRIMARY KEY (id)  
);
```

Second Schema Migration

```
ALTER TABLE  
Users  
ADD COLUMN  
age_int int;  
  
ALTER TABLE  
Users  
ADD COLUMN  
status boolean
```

Index Migration

```
CREATE  
INDEX  
CONCURRENTLY  
status_email  
ON Users  
(email, status)
```

Schema

```
CREATE TABLE  
Users(  
    id bigint  
    first_name varchar(),  
    last_name varchar(),  
    email varchar(),  
    phone varchar(),  
    PRIMARY KEY (id)  
);
```

Second Schema Migration

```
ALTER TABLE  
Users  
ADD COLUMN  
age_int int;  
  
ALTER TABLE  
Users  
ADD COLUMN  
status boolean
```

Index Migration

```
CREATE  
INDEX  
CONCURRENTLY  
status_email  
ON Users  
(email, status)
```

Third Schema Migration

```
ALTER TABLE  
Users  
ADD COLUMN  
active_date  
datetime;
```

Schema

```
CREATE TABLE  
Users(  
    id bigint  
    first_name varchar(),  
    last_name varchar(),  
    email varchar(),  
    phone varchar(),  
    PRIMARY KEY (id)  
);
```

Second Schema Migration

```
ALTER TABLE  
Users  
ADD COLUMN  
age_int int;  
  
ALTER TABLE  
Users  
ADD COLUMN  
status boolean
```

Index Migration

```
CREATE  
INDEX  
CONCURRENTLY  
status_email  
ON Users  
(email, status)
```

Third Schema Migration

```
ALTER TABLE  
Users  
ADD COLUMN  
active_date  
datetime;
```

Data Migration

```
UPDATE  
Users  
SET  
active_date=now()  
WHERE  
status
```

Back to Locking in theory

Schema

```
CREATE TABLE  
Users(  
    id bigint  
    first_name varchar(),  
    last_name varchar(),  
    email varchar(),  
    phone varchar(),  
    PRIMARY KEY (id)  
);
```

Second Schema Migration

```
ALTER TABLE  
Users  
ADD COLUMN  
age_int int;  
  
ALTER TABLE  
Users  
ADD COLUMN  
status boolean
```

ACCESS
EXCLUSIVE

Index Migration

```
CREATE  
INDEX  
CONCURRENTLY  
status_email  
ON Users  
(email, status)
```

Third Schema Migration

```
ALTER TABLE  
Users  
ADD COLUMN  
active_date  
datetime;
```

Data Migration

```
UPDATE  
Users  
SET  
active_date=now()  
WHERE  
status
```

Second Schema Migration

```
ALTER TABLE  
Users  
ADD COLUMN  
age_int int;
```

```
ALTER TABLE  
Users  
ADD COLUMN  
status boolean
```

ACCESS
EXCLUSIVE

SELECT *
FROM USERS

Third Schema Migration

```
ALTER TABLE  
Users  
ADD COLUMN  
active_date  
datetime;
```

ACCESS
EXCLUSIVE

ways to get around locking

Changing the type of a column

Add a new column with new type (**Schema**)

Write to both columns and then backfill new column. (**Data**)

New version of the code stops writing to the old column. (**Process**)

ways to get around locking

Add a column that has a default:

Add column (**SCHEMA**)

Add default as a separate command (**DATA**)

Backfill the column with the default value (**DATA**)

scaling with tools

Don't lock yourself
out

Break your
migrations

scaling with tools

Don't lock yourself
out

Break your
migrations

Database +
Code Compatibility

database + code compatibility

- Backwards AND forward compatible
- Dropping columns (future)

scaling with tools

Don't lock yourself
out

Break your
migrations

Database +
Code Compatibility

Migration User

scaling with your
team

migration - user account

Make a user to run your migrations

```
ALTER ROLE jumana  
SET lock_timeout='5s' ;
```

USERS TABLE

SELECT

SELECT

SELECT

SELECT

SELECT

USERS TABLE

ALTER

SELECT

SELECT

SELECT

SELECT

SELECT

USERS TABLE

SELECT

SELECT

SELECT

ALTER

SELECT

SELECT

SELECT

SELECT

SELECT

USERS TABLE

USERS TABLE

SELECT

SELECT

SELECT

ALTER

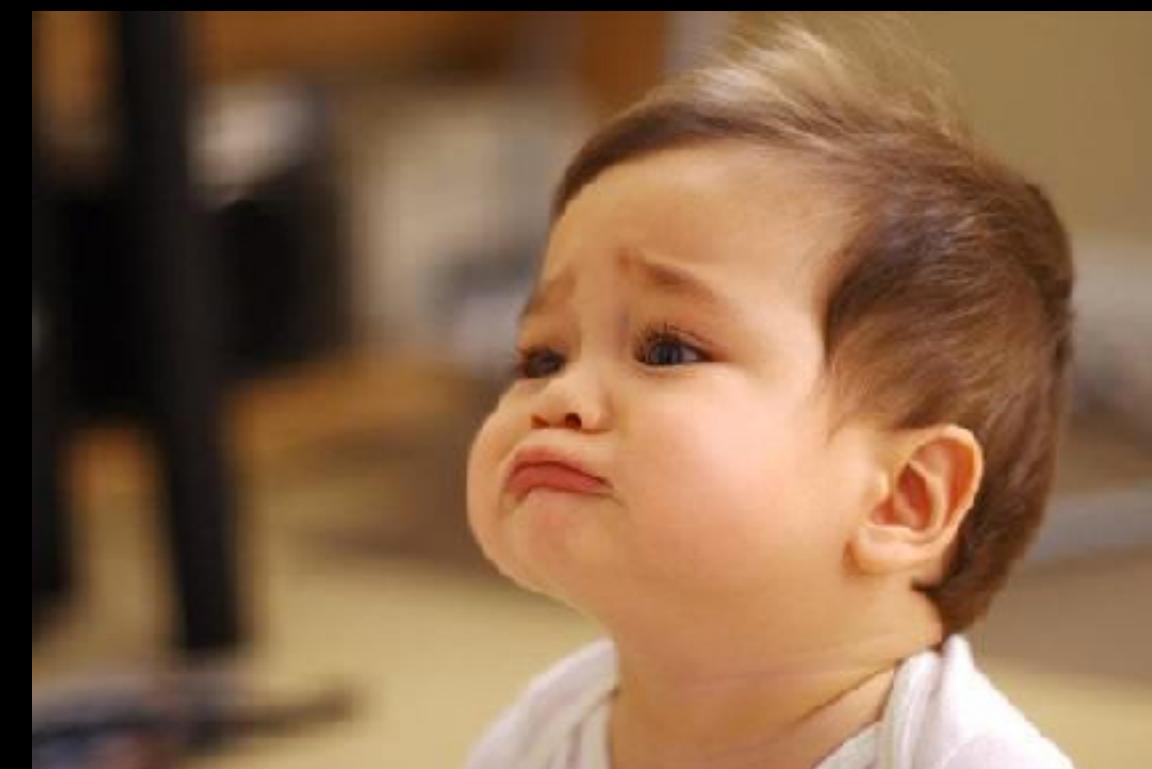
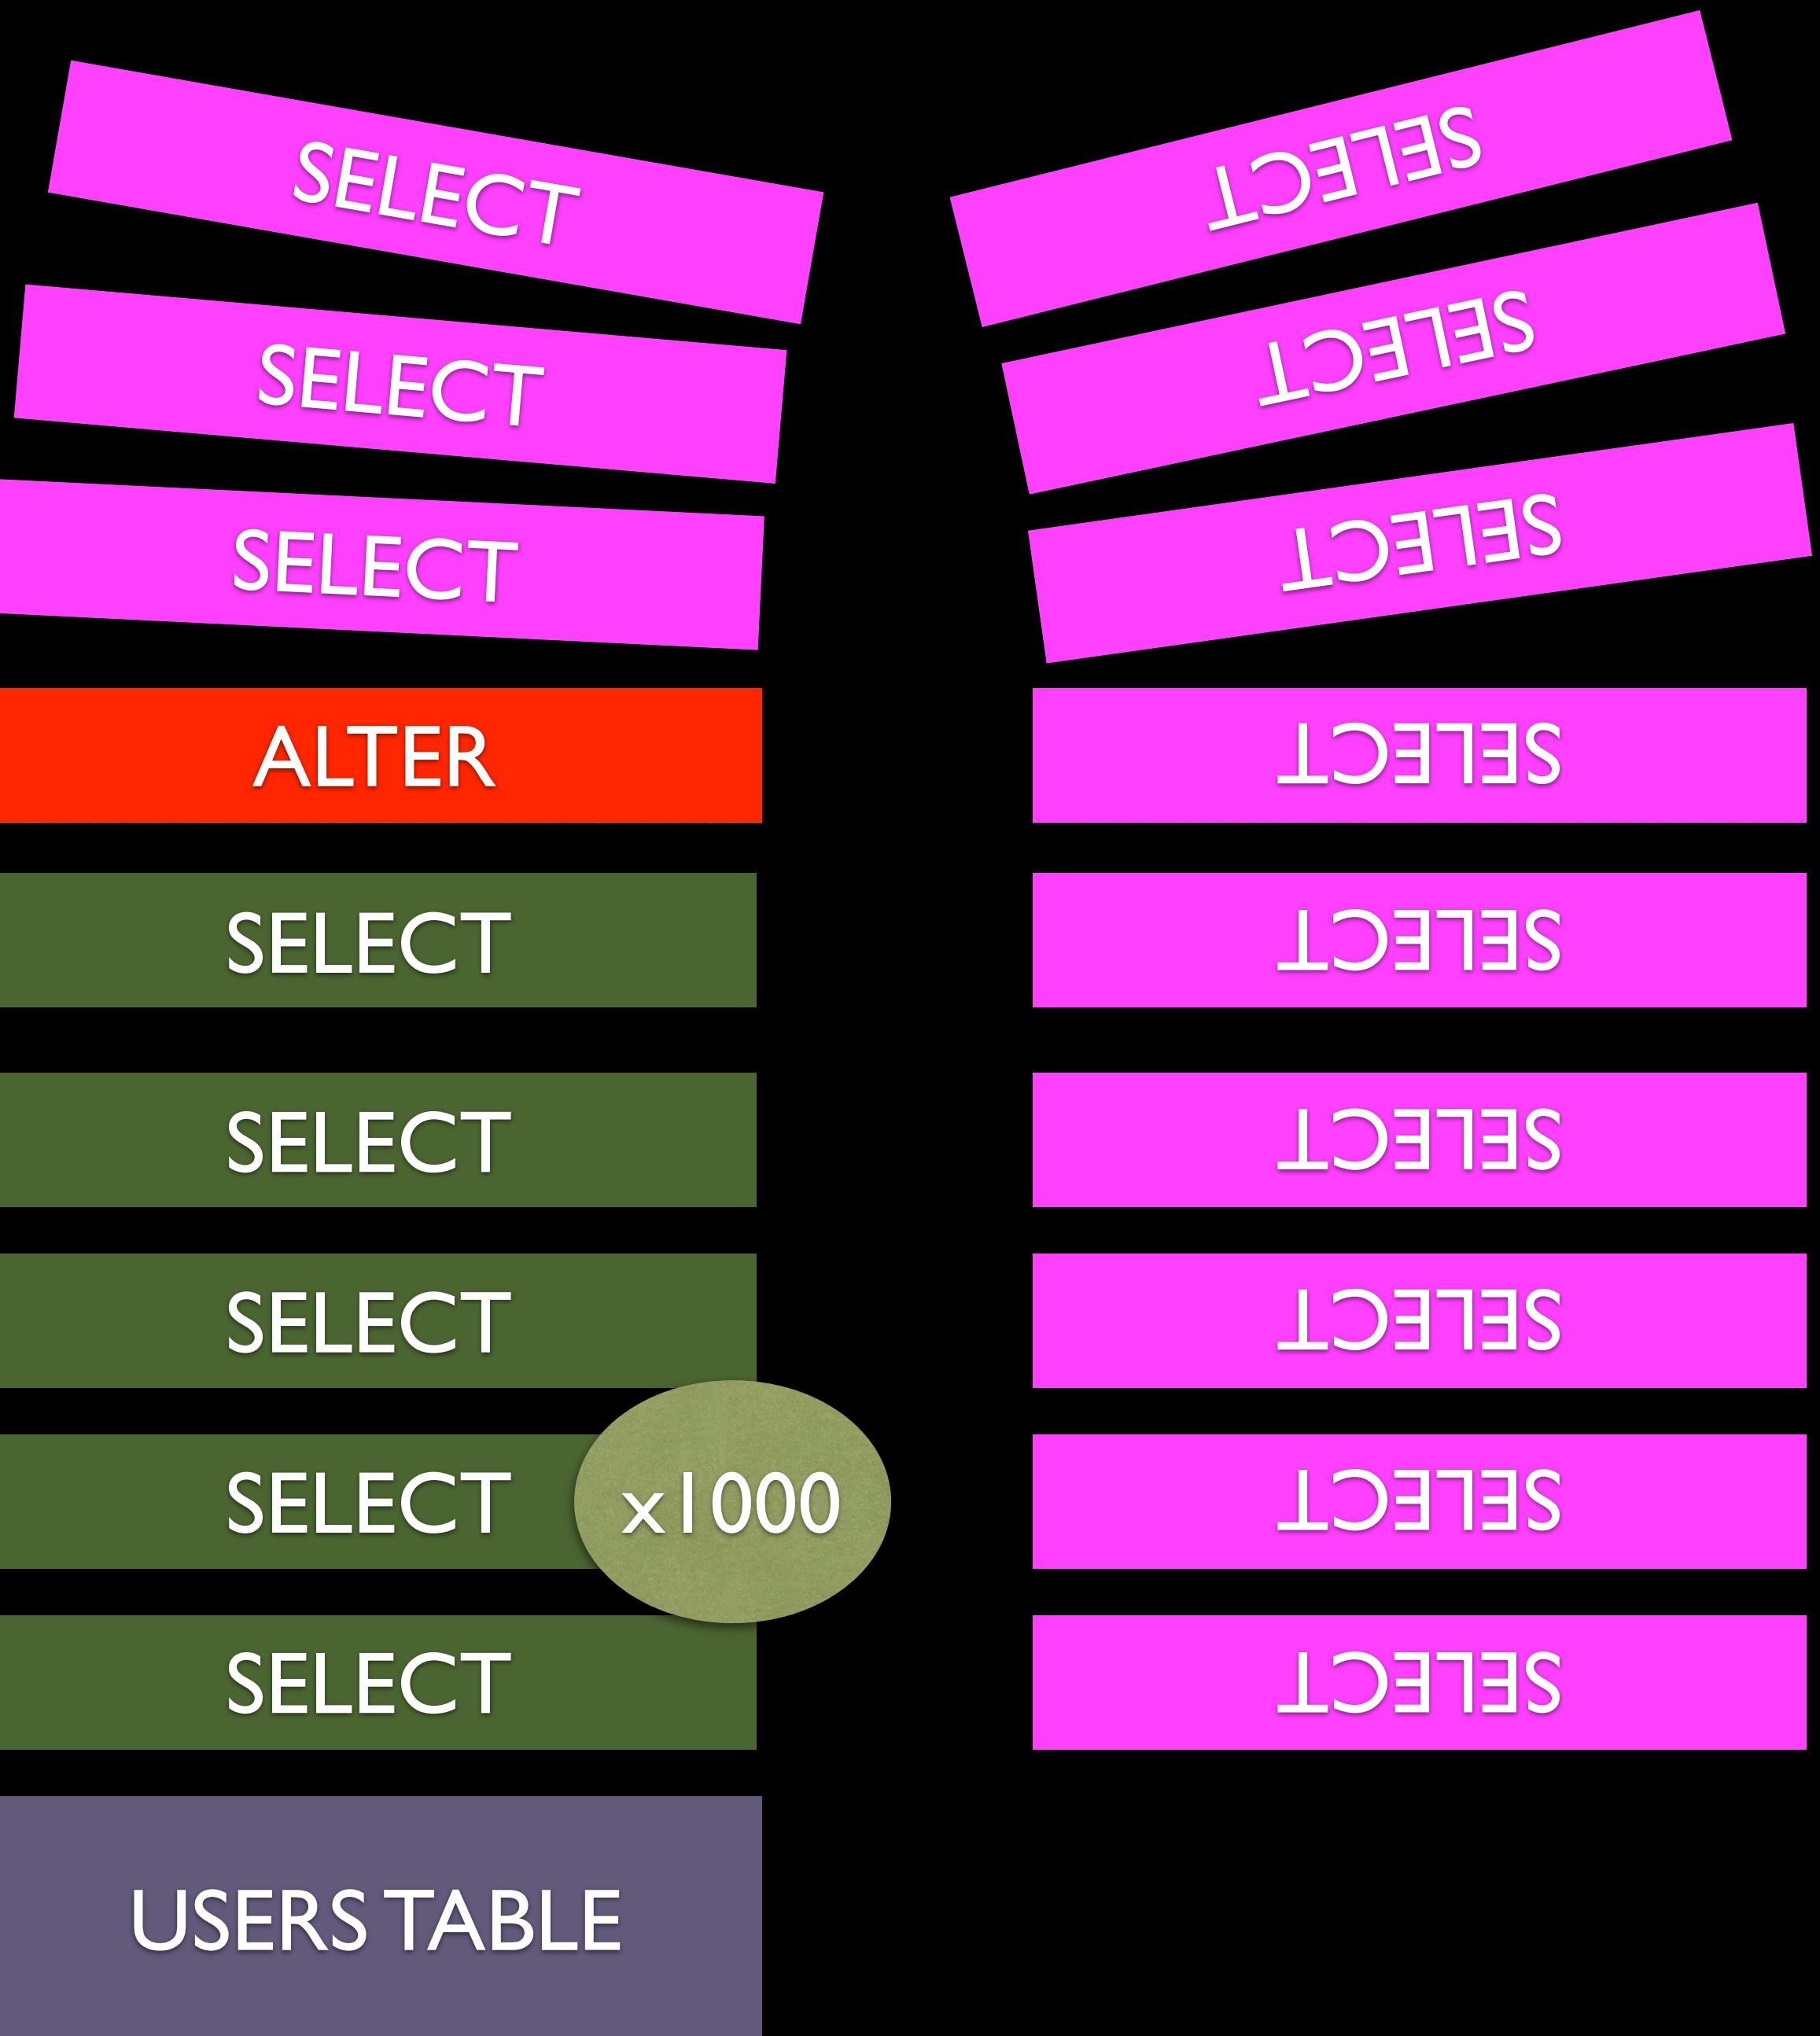
SELECT

SELECT

SELECT

SELECT x1000

SELECT

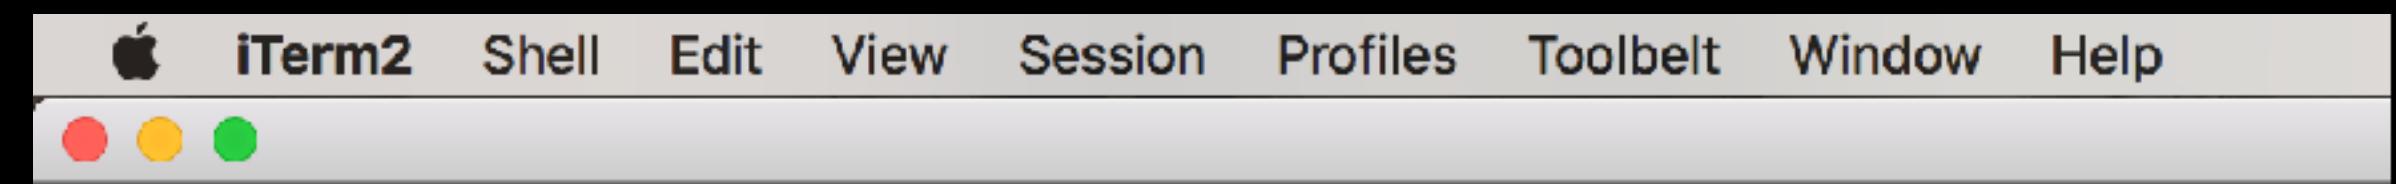


scaling with your
team

Migration User

Migration Tooling

migration tooling



~ psql -U jz -h localhost

jz=> create table users



migration tooling

version_num

39b980e419a9

7ac4c29288a3

16af97dd95c1

24187609b002

9d3ef281fb30

8918b762b17b

47f9fc04d6a7

2badb70854a9

1551e5d79984

Automate Scripts

migration tooling

version_num

39b980e419a9

7ac4c29288a3

16af97dd95c1

24187609b002

9d3ef281fb30

8918b762b17b

47f9fc04d6a7

2badb70854a9

1551e5d79984

Automate Scripts

State of your DB

migration tooling

version_num

39b980e419a9

7ac4c29288a3

16af97dd95c1

24187609b002

9d3ef281fb30

8918b762b17b

47f9fc04d6a7

2badb70854a9

1551e5d79984

Automate Scripts

State of your DB

Easy replication on local
dev

scaling with your
team

Migration User

Migration Tooling
Load Testing

load test your migrations

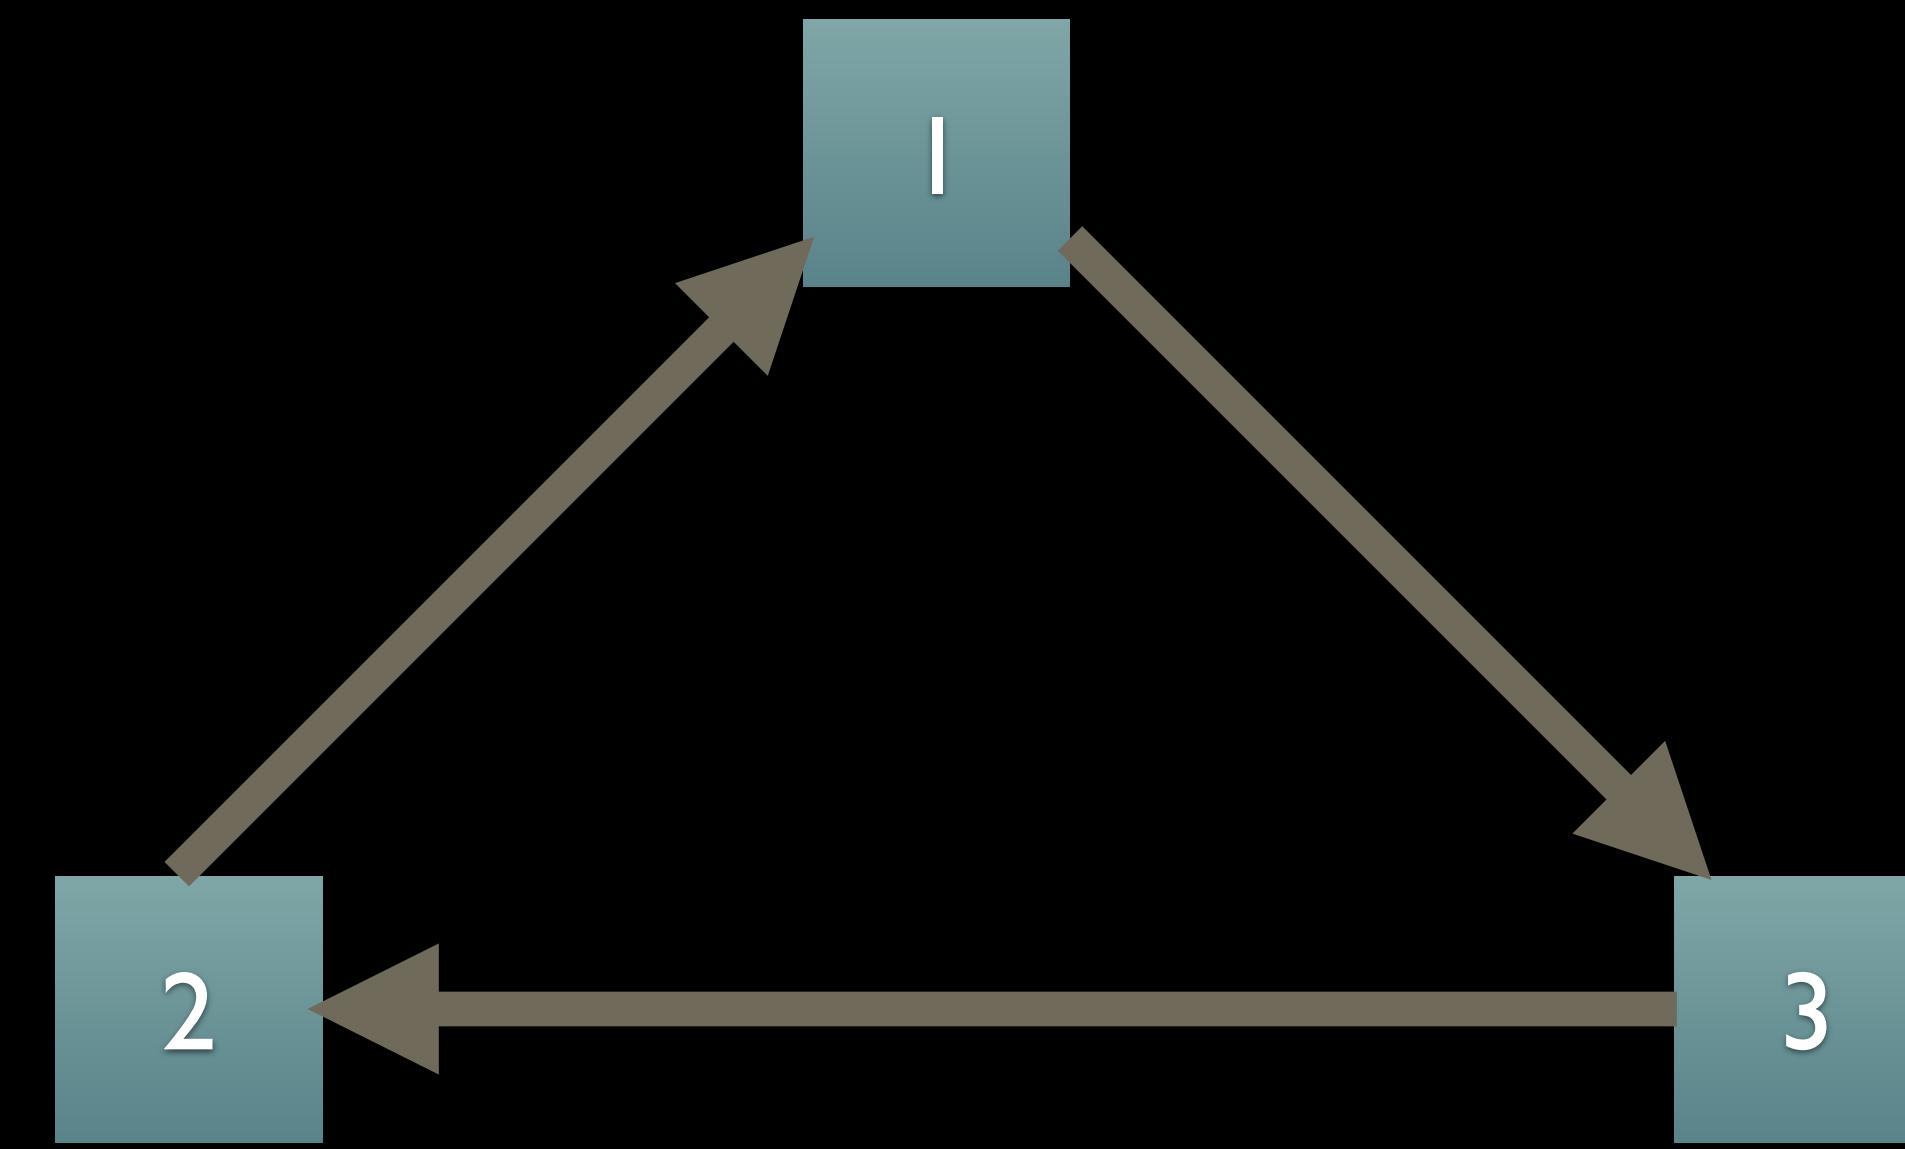
- Run migrations during low traffic

load test your migrations

- Run migrations during low traffic
- What if low traffic still means you have 1000's of users on your site?

load test your migrations

- Run migrations during low traffic
- What if low traffic still means you have 1000's of users on your site?
- Load test your migration



TIME

HTTP REQUEST CYCLE

MIGRATION

TIME

1

HTTP REQUEST CYCLE

Request cycle start (**Transaction 1**)
Update Users set status=False where id=2;

MIGRATION

TIME

1

HTTP REQUEST CYCLE
Request cycle start (**Transaction 1**)
Update Users set status=False where id=2;
Spawn async workers

MIGRATION

TIME

1

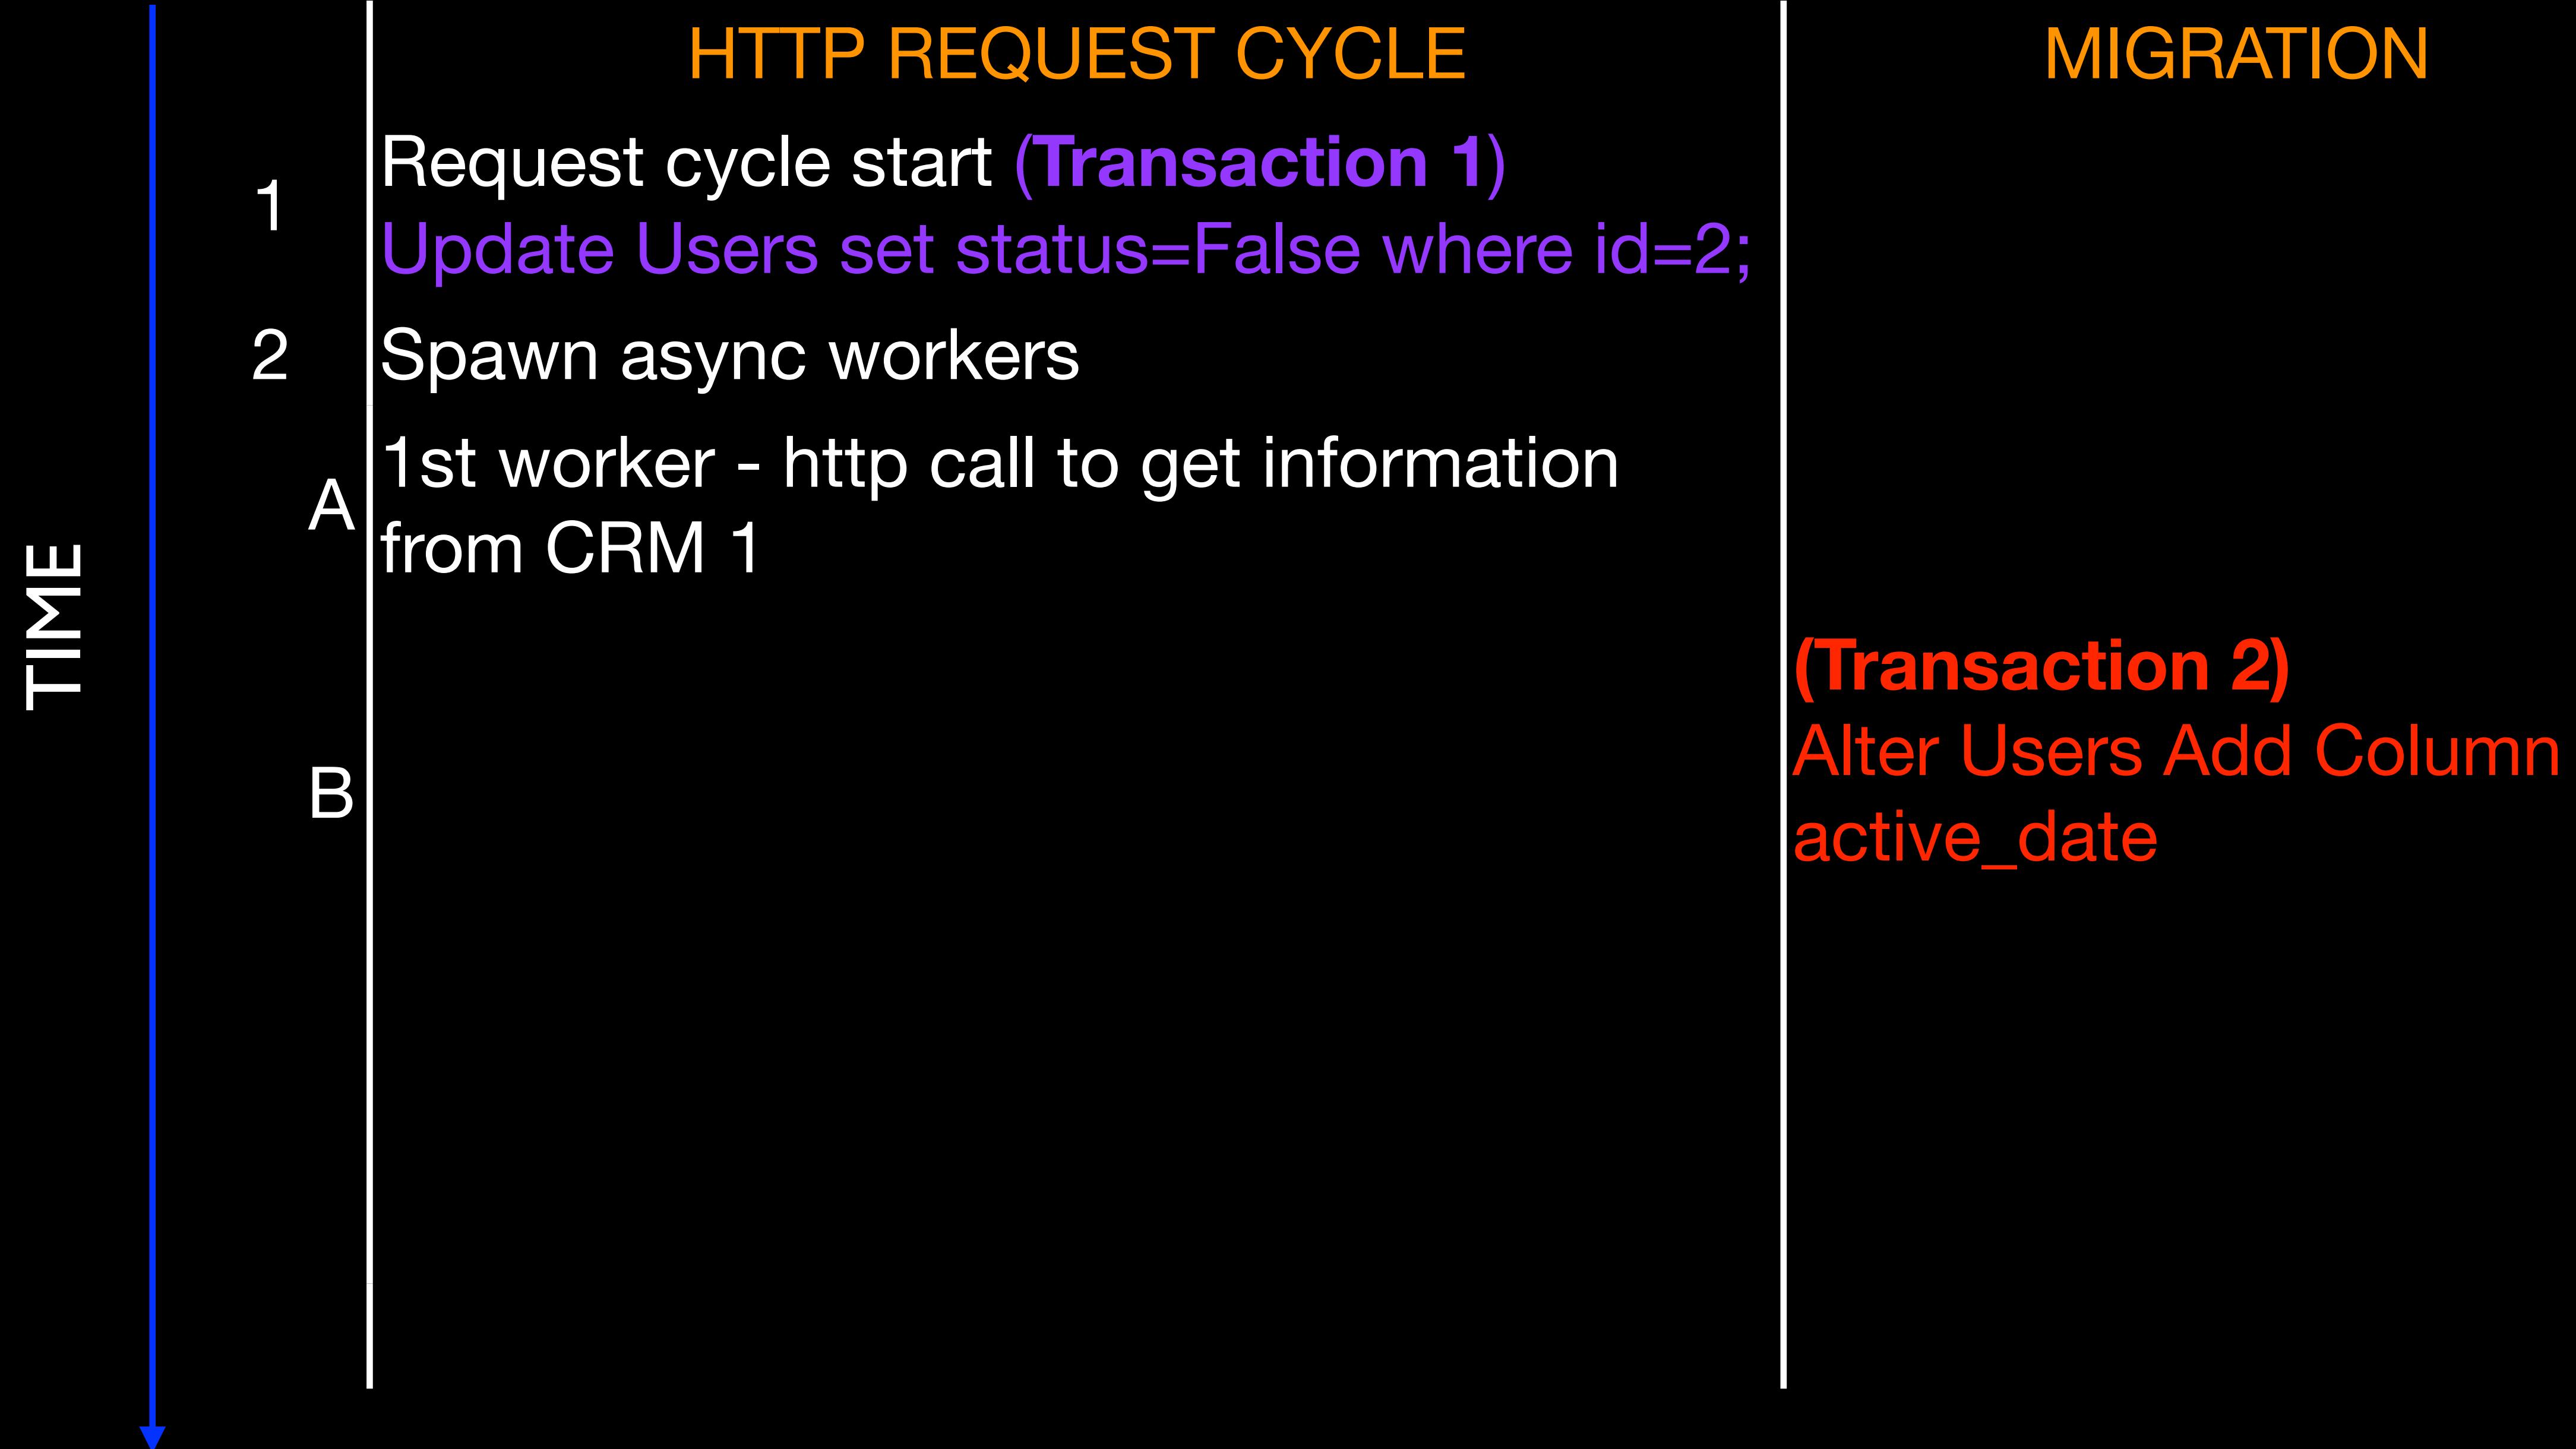
HTTP REQUEST CYCLE

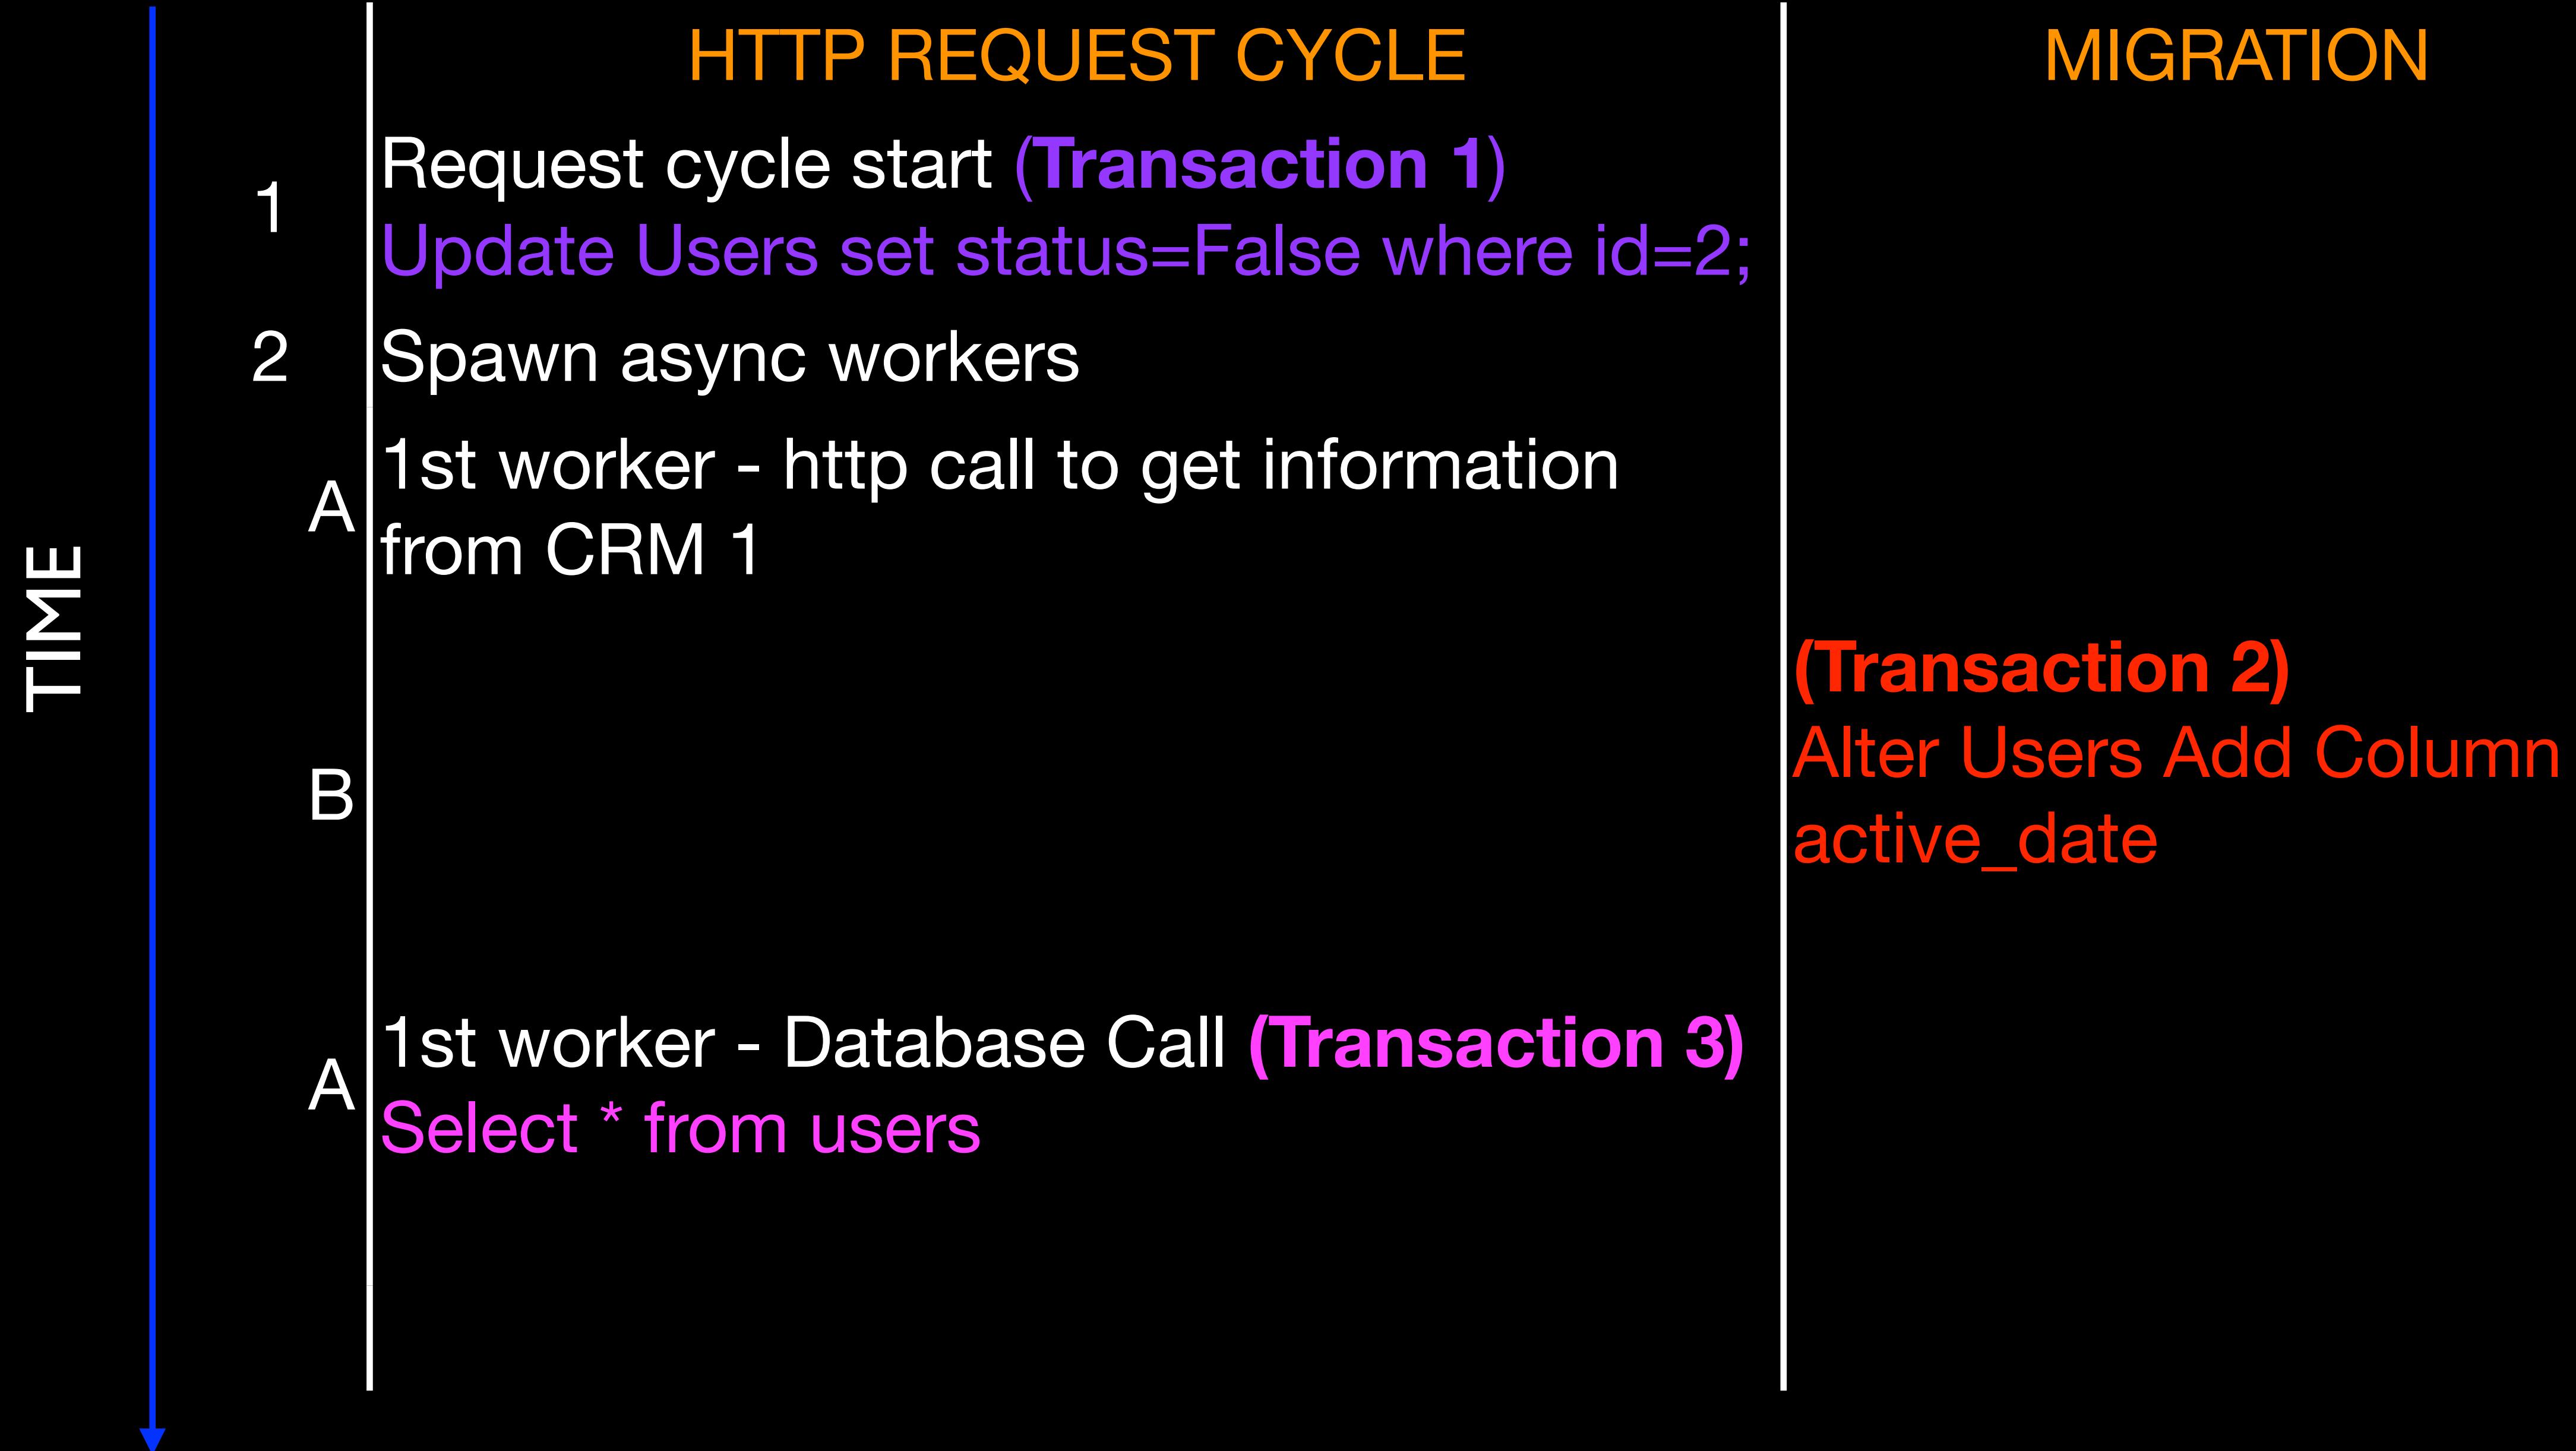
Request cycle start (**Transaction 1**)
Update Users set status=False where id=2;
Spawn async workers
1st worker - http call to get information
from CRM 1

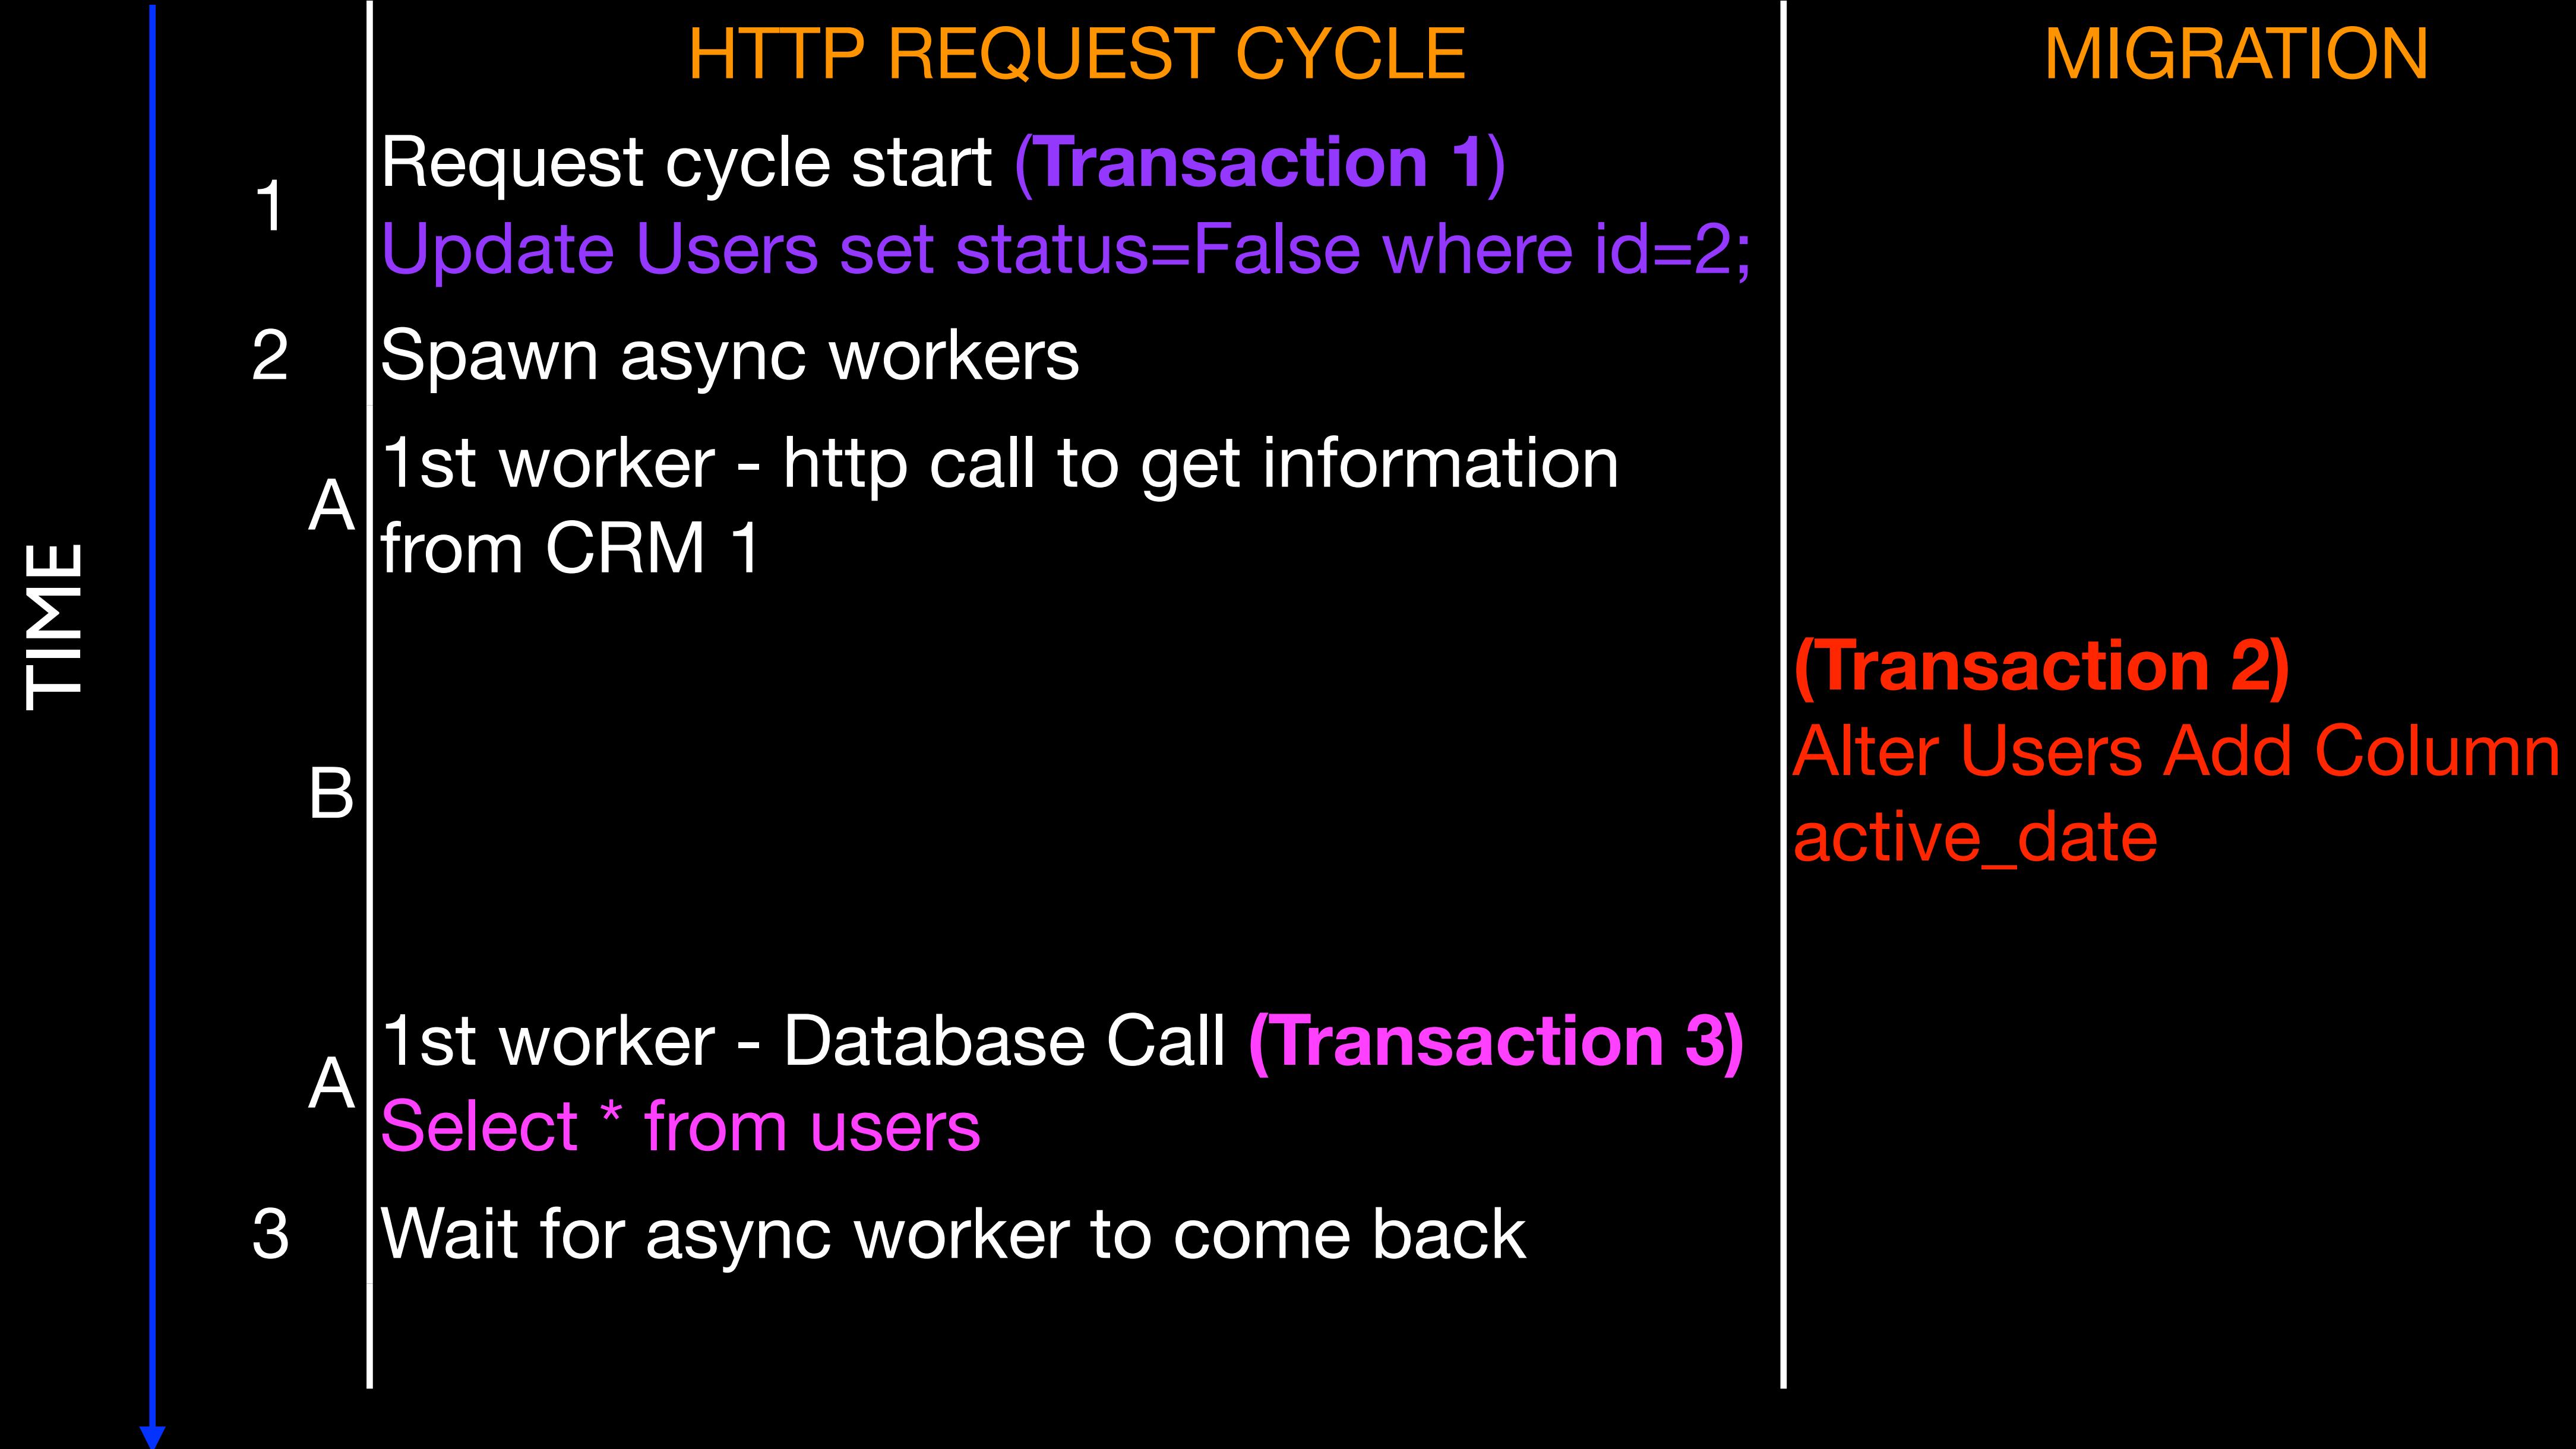
2

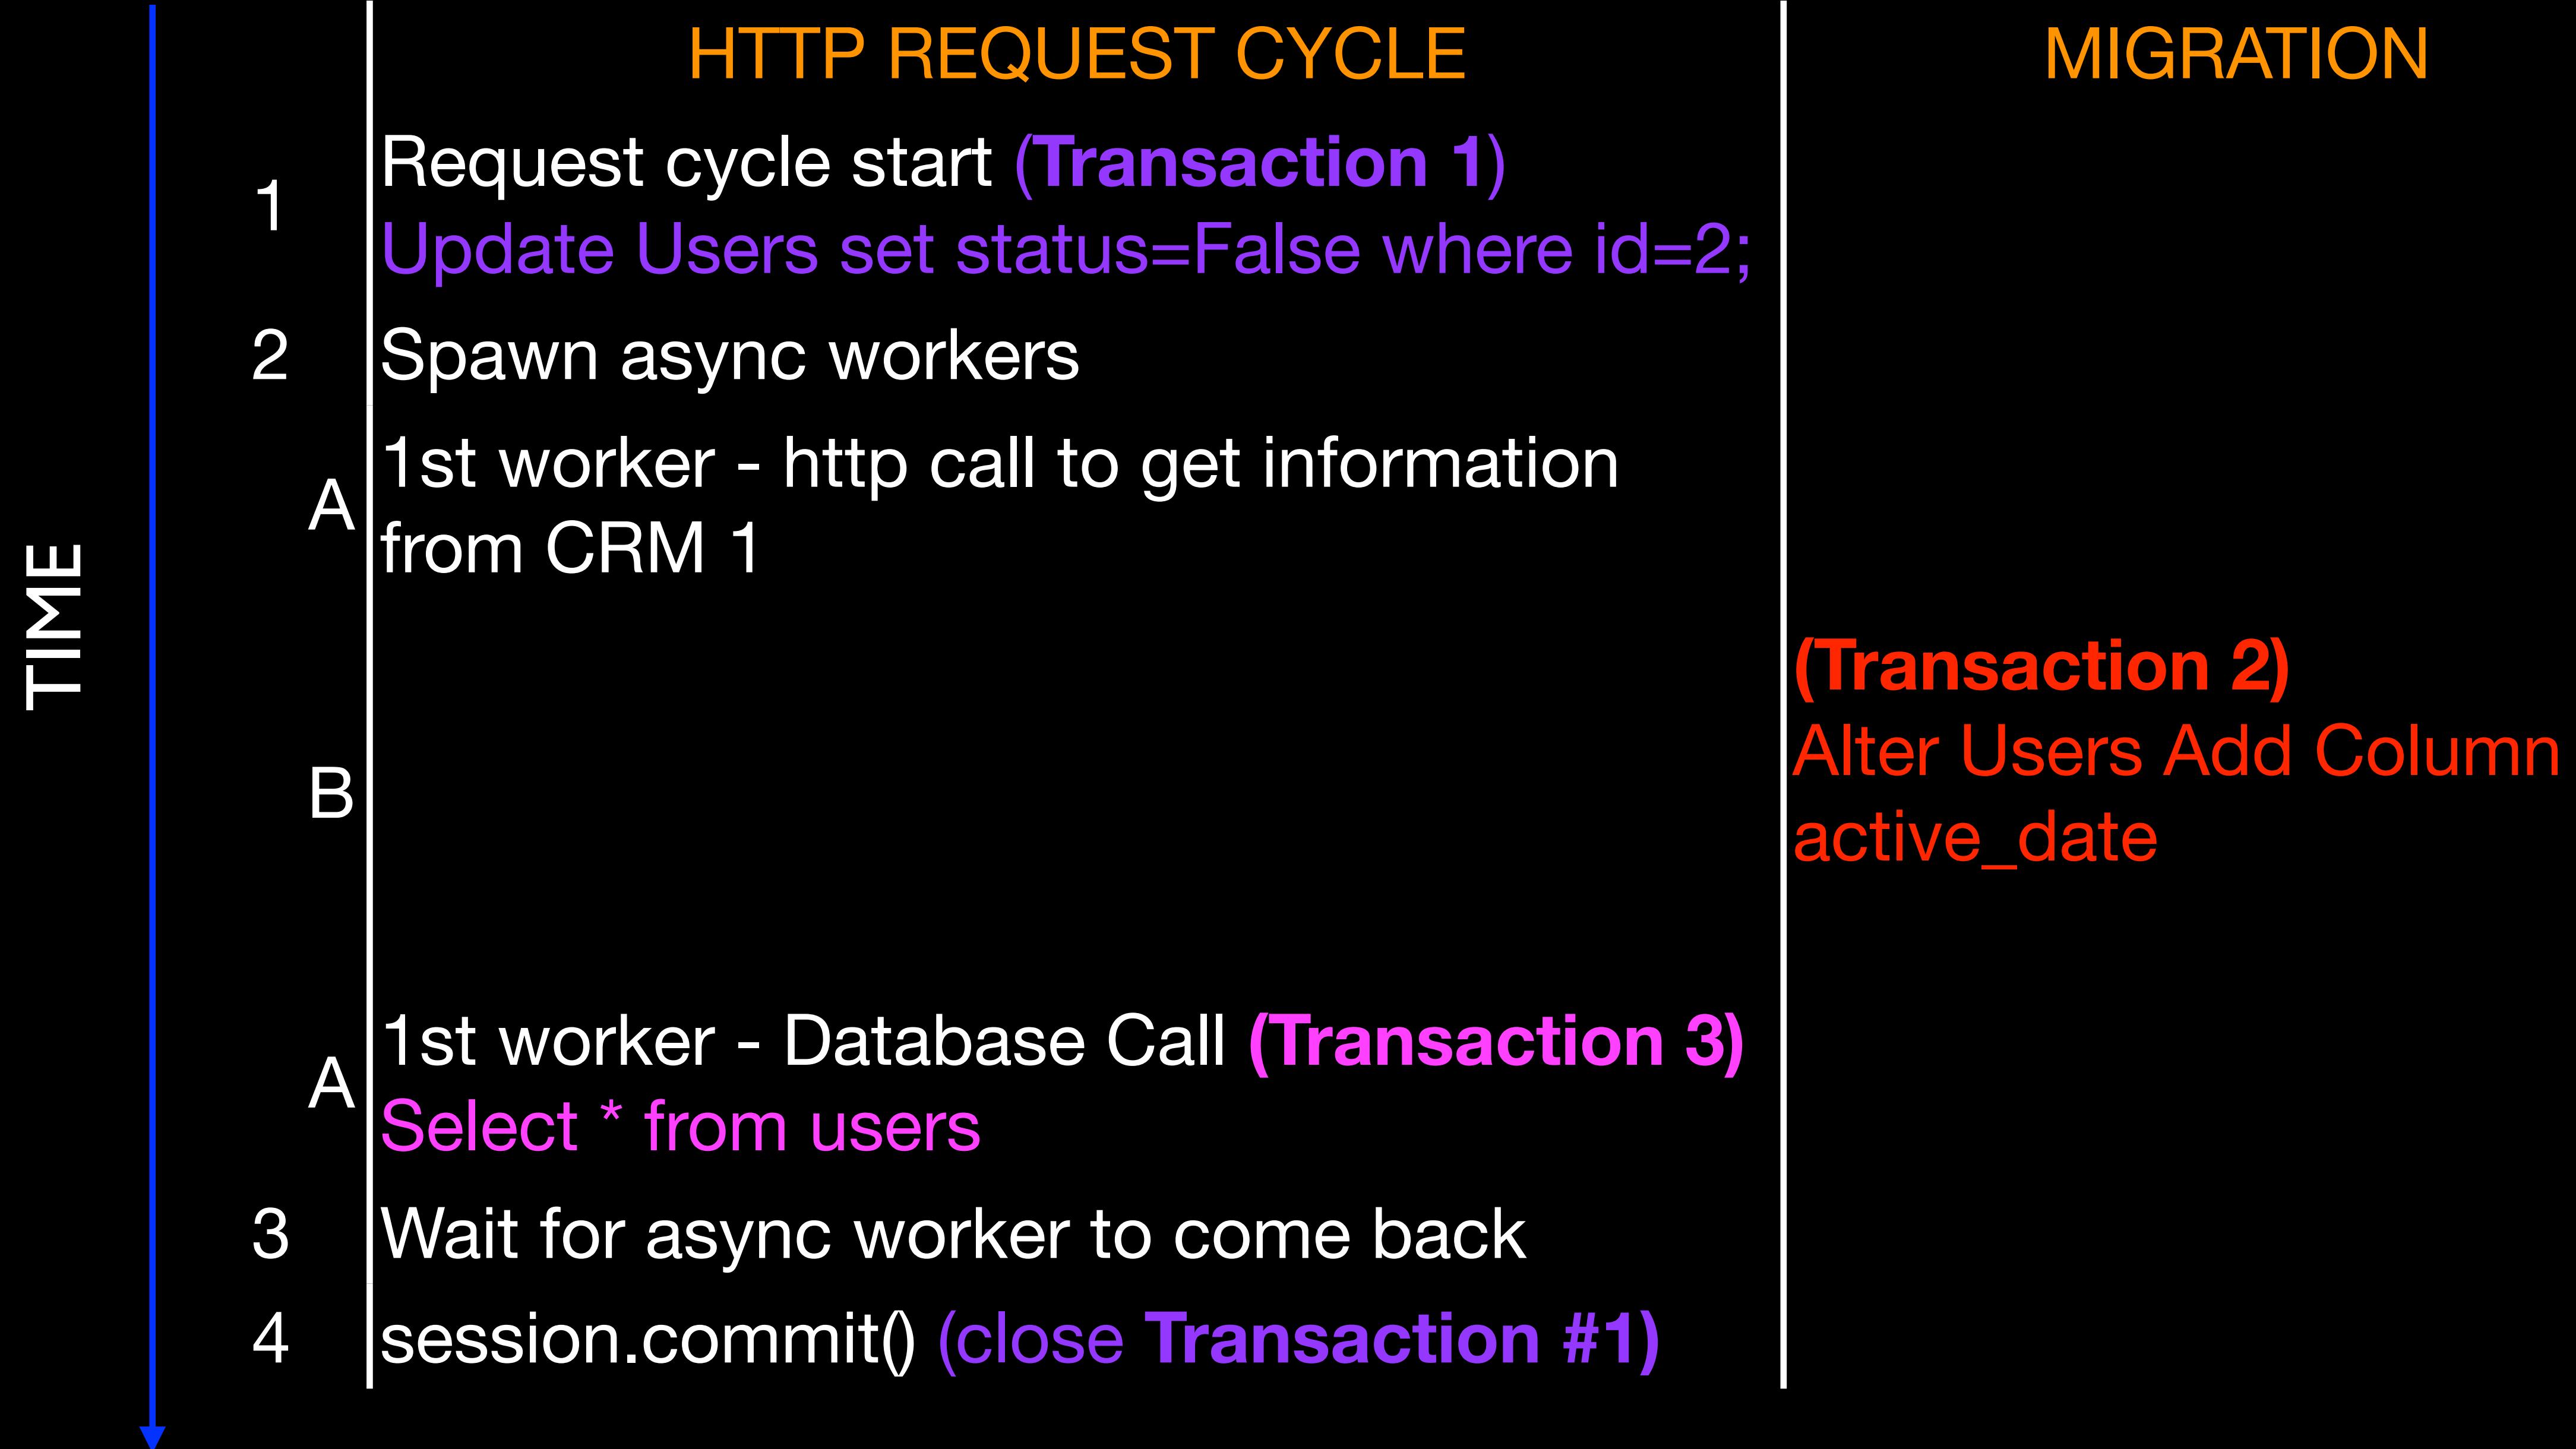
A

MIGRATION









TRANSACTION I

HTTP CALL,

REQUEST CYCLE,

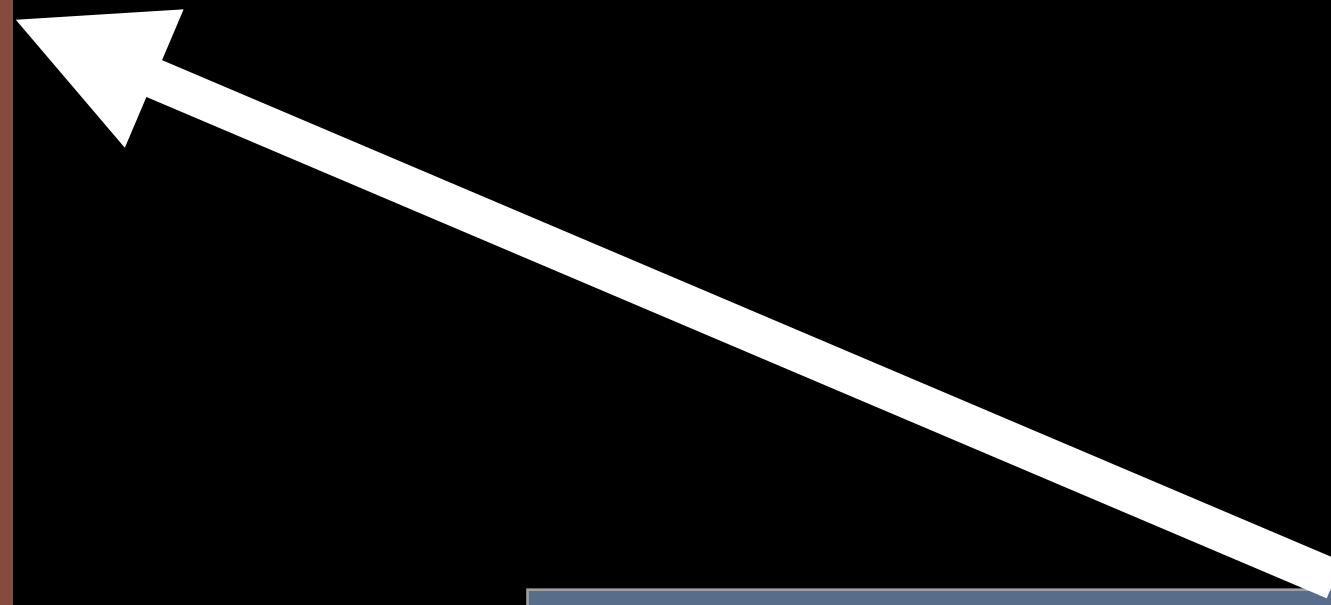
Update Users set status=False
where id=2;

TRANSACTION 1

HTTP CALL,
REQUEST CYCLE,
Update Users set status=False
where id=2;

TRANSACTION 2

MIGRATION
Alter Users Add Column
active_date



TRANSACTION 1

HTTP CALL,

REQUEST CYCLE,

Update Users set status=False
where id=2;

TRANSACTION 2

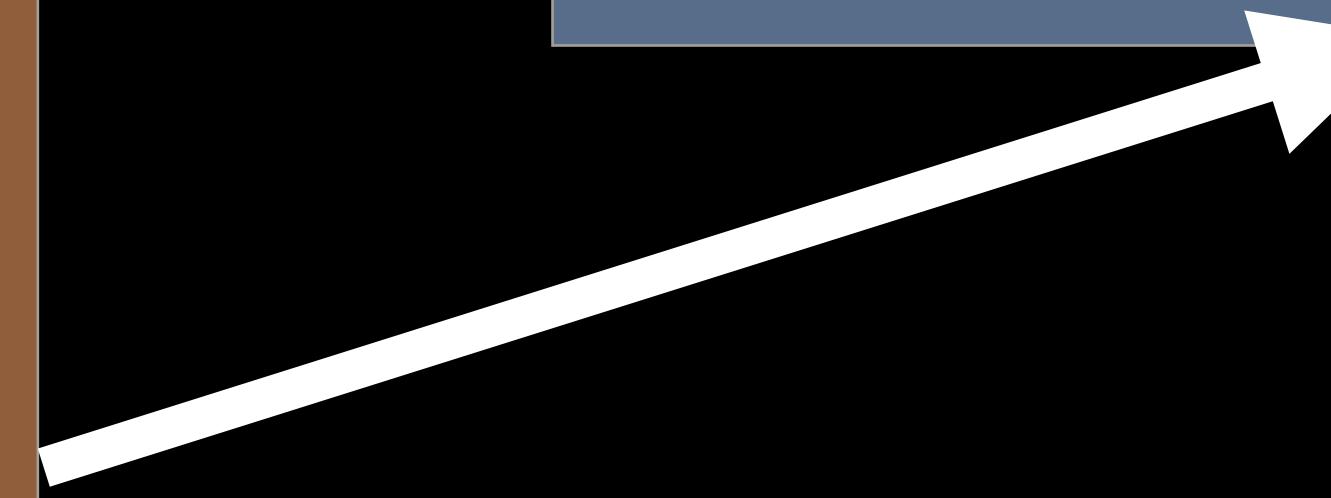
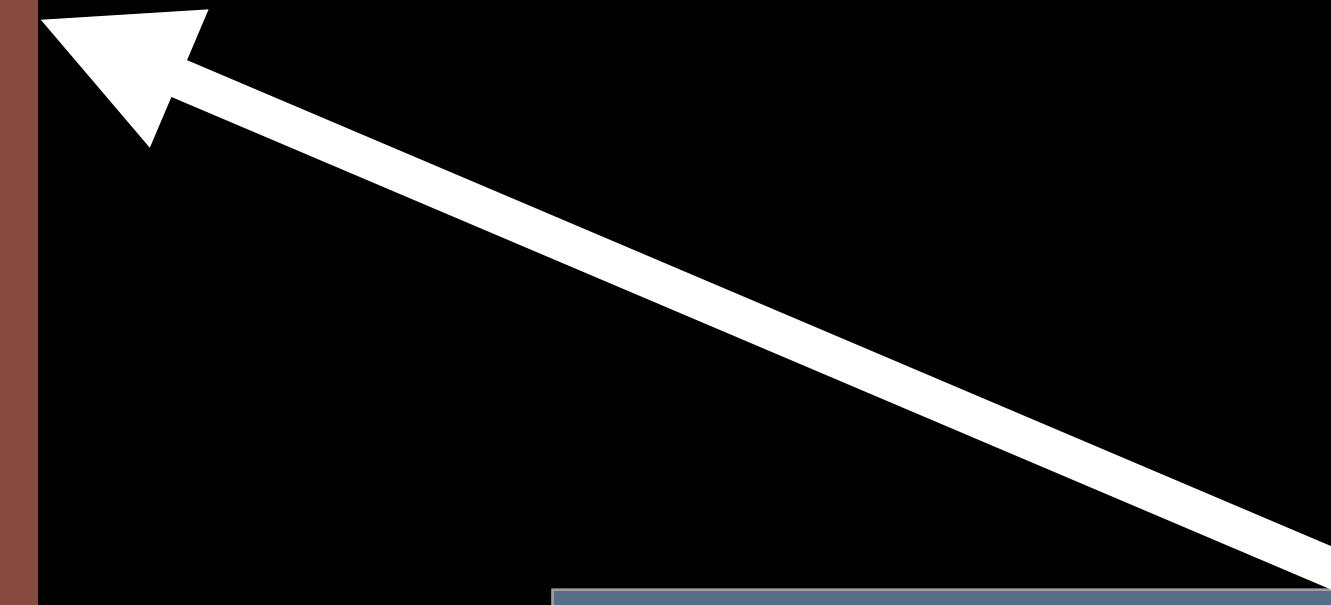
MIGRATION

Alter Users Add Column
active_date

TRANSACTION 3

ASYNC WORKER

Select * from users



TRANSACTION 1

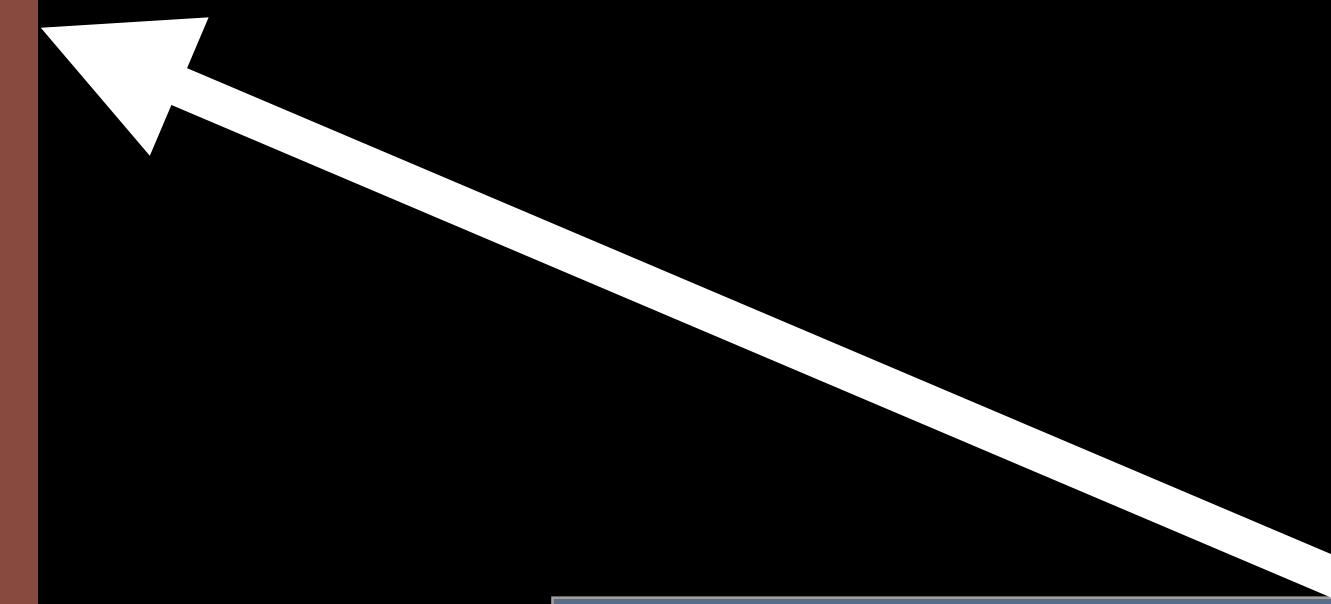
HTTP CALL,
REQUEST CYCLE,
Update Users set status=False
where id=2;

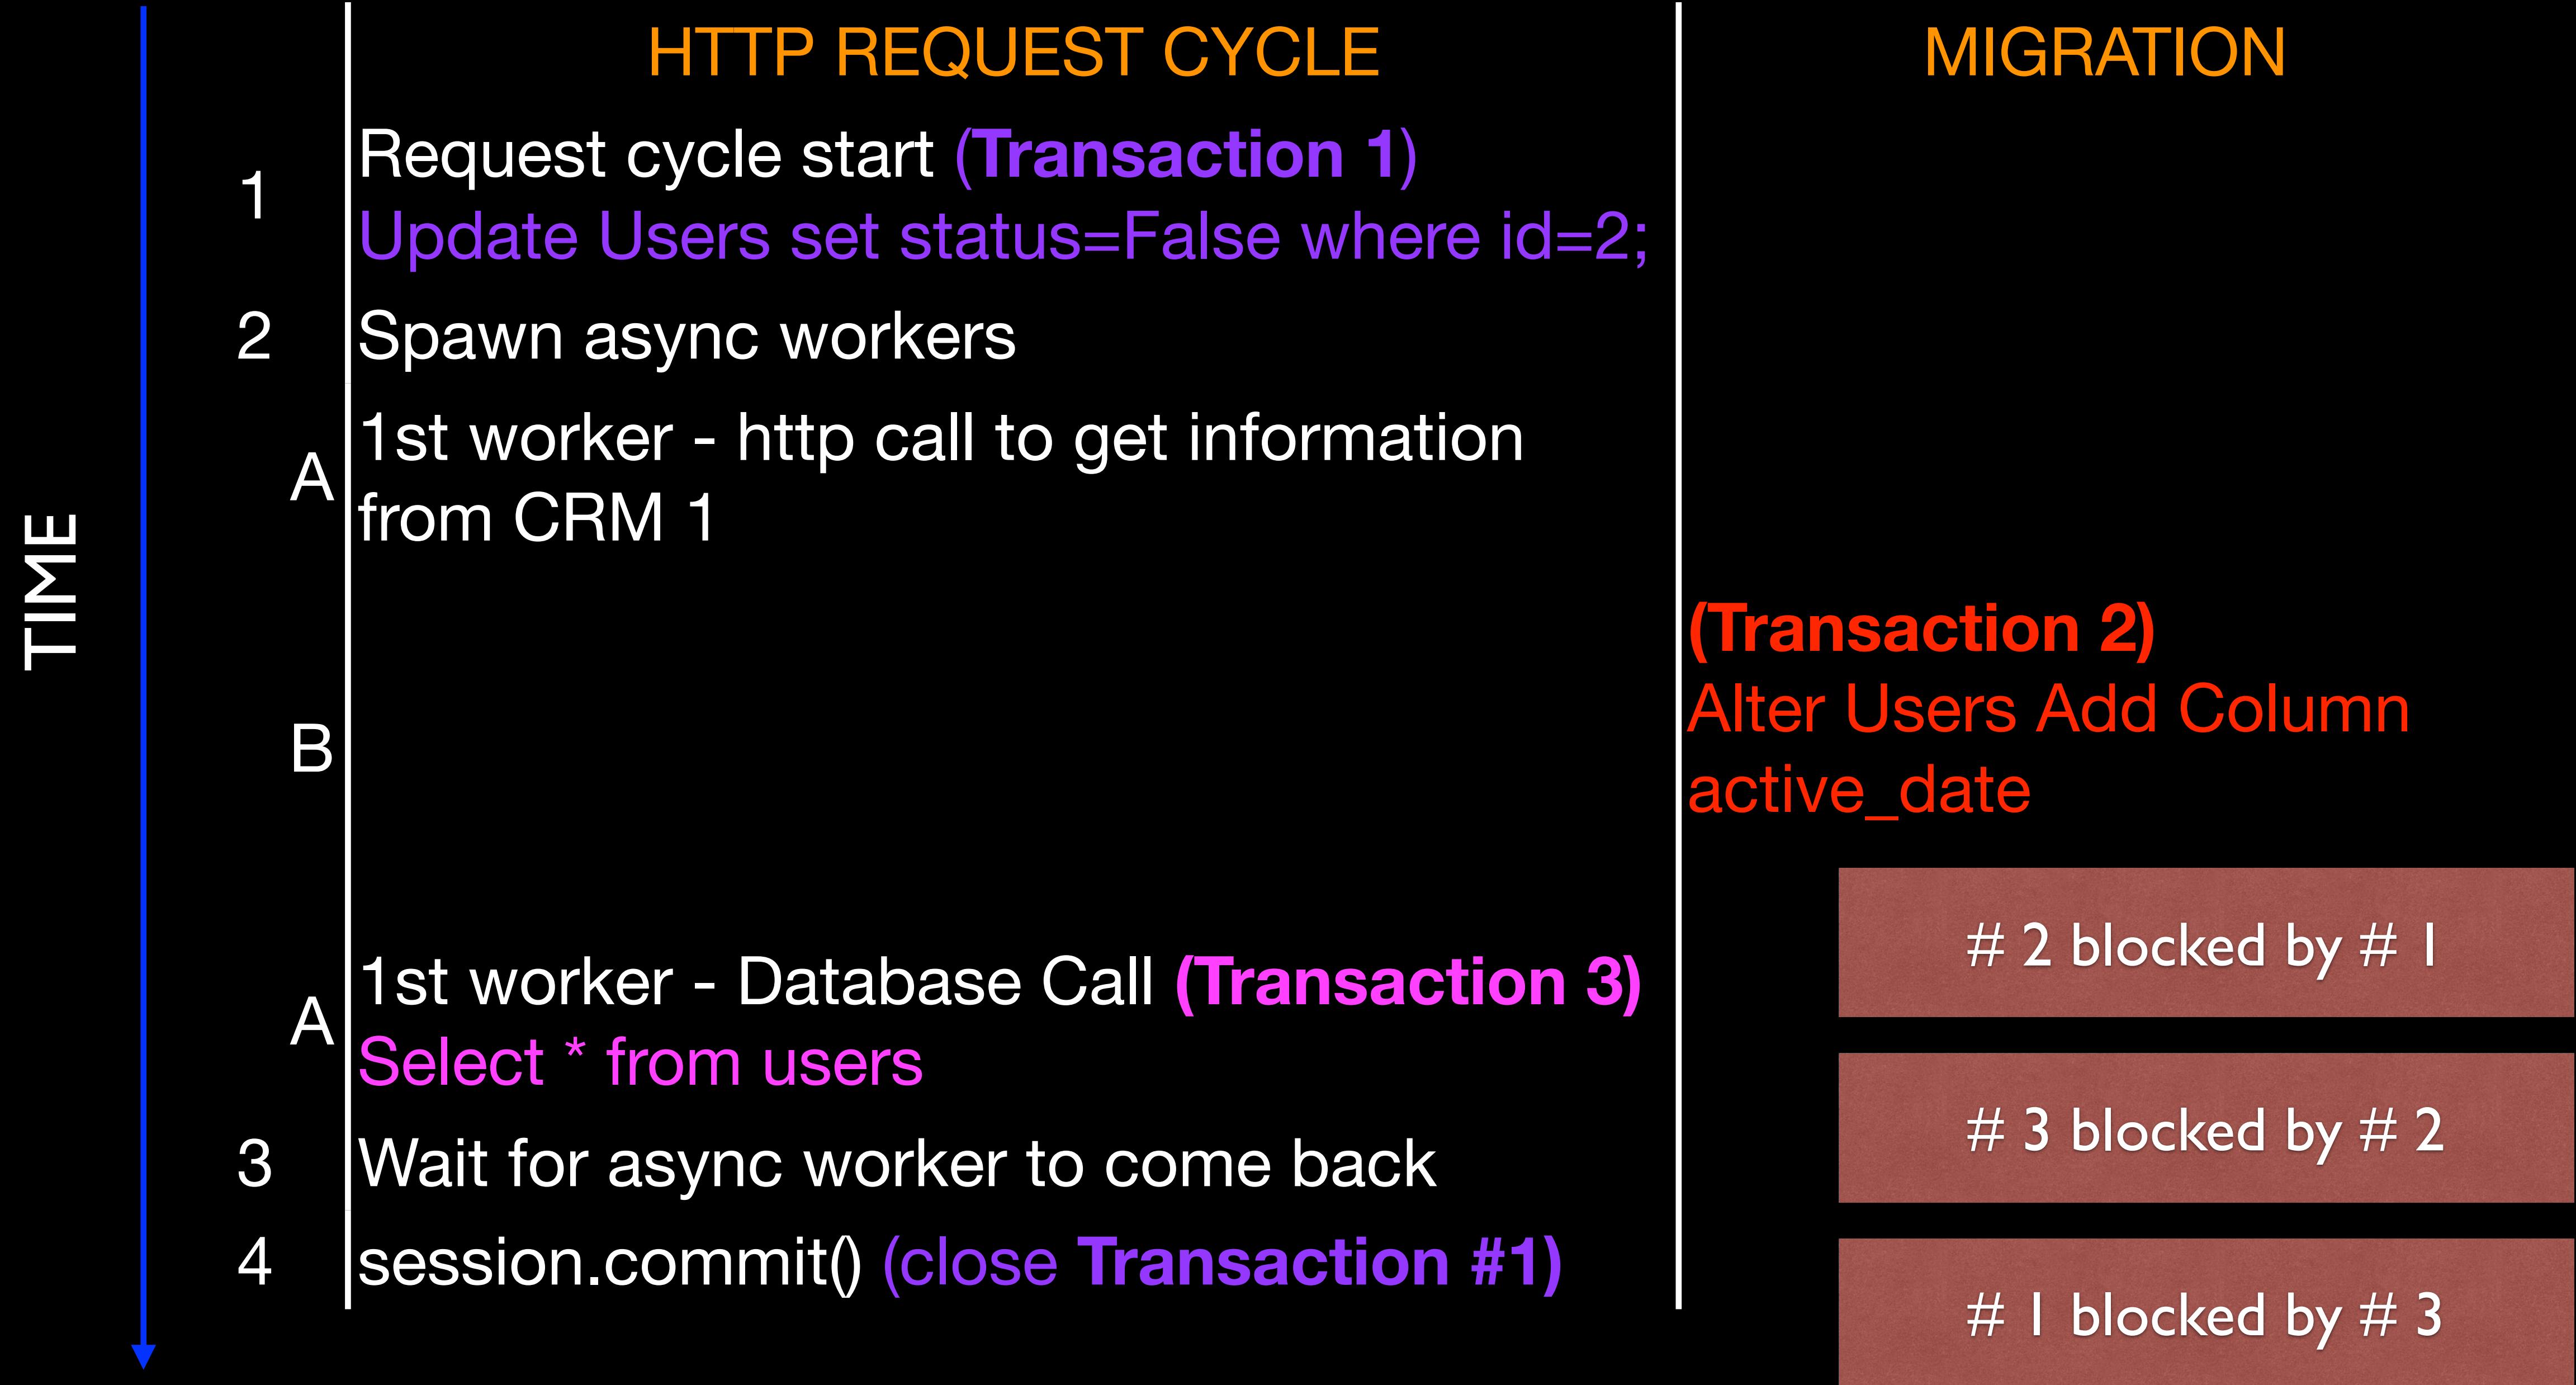
TRANSACTION 2

MIGRATION
Alter Users Add Column
active_date

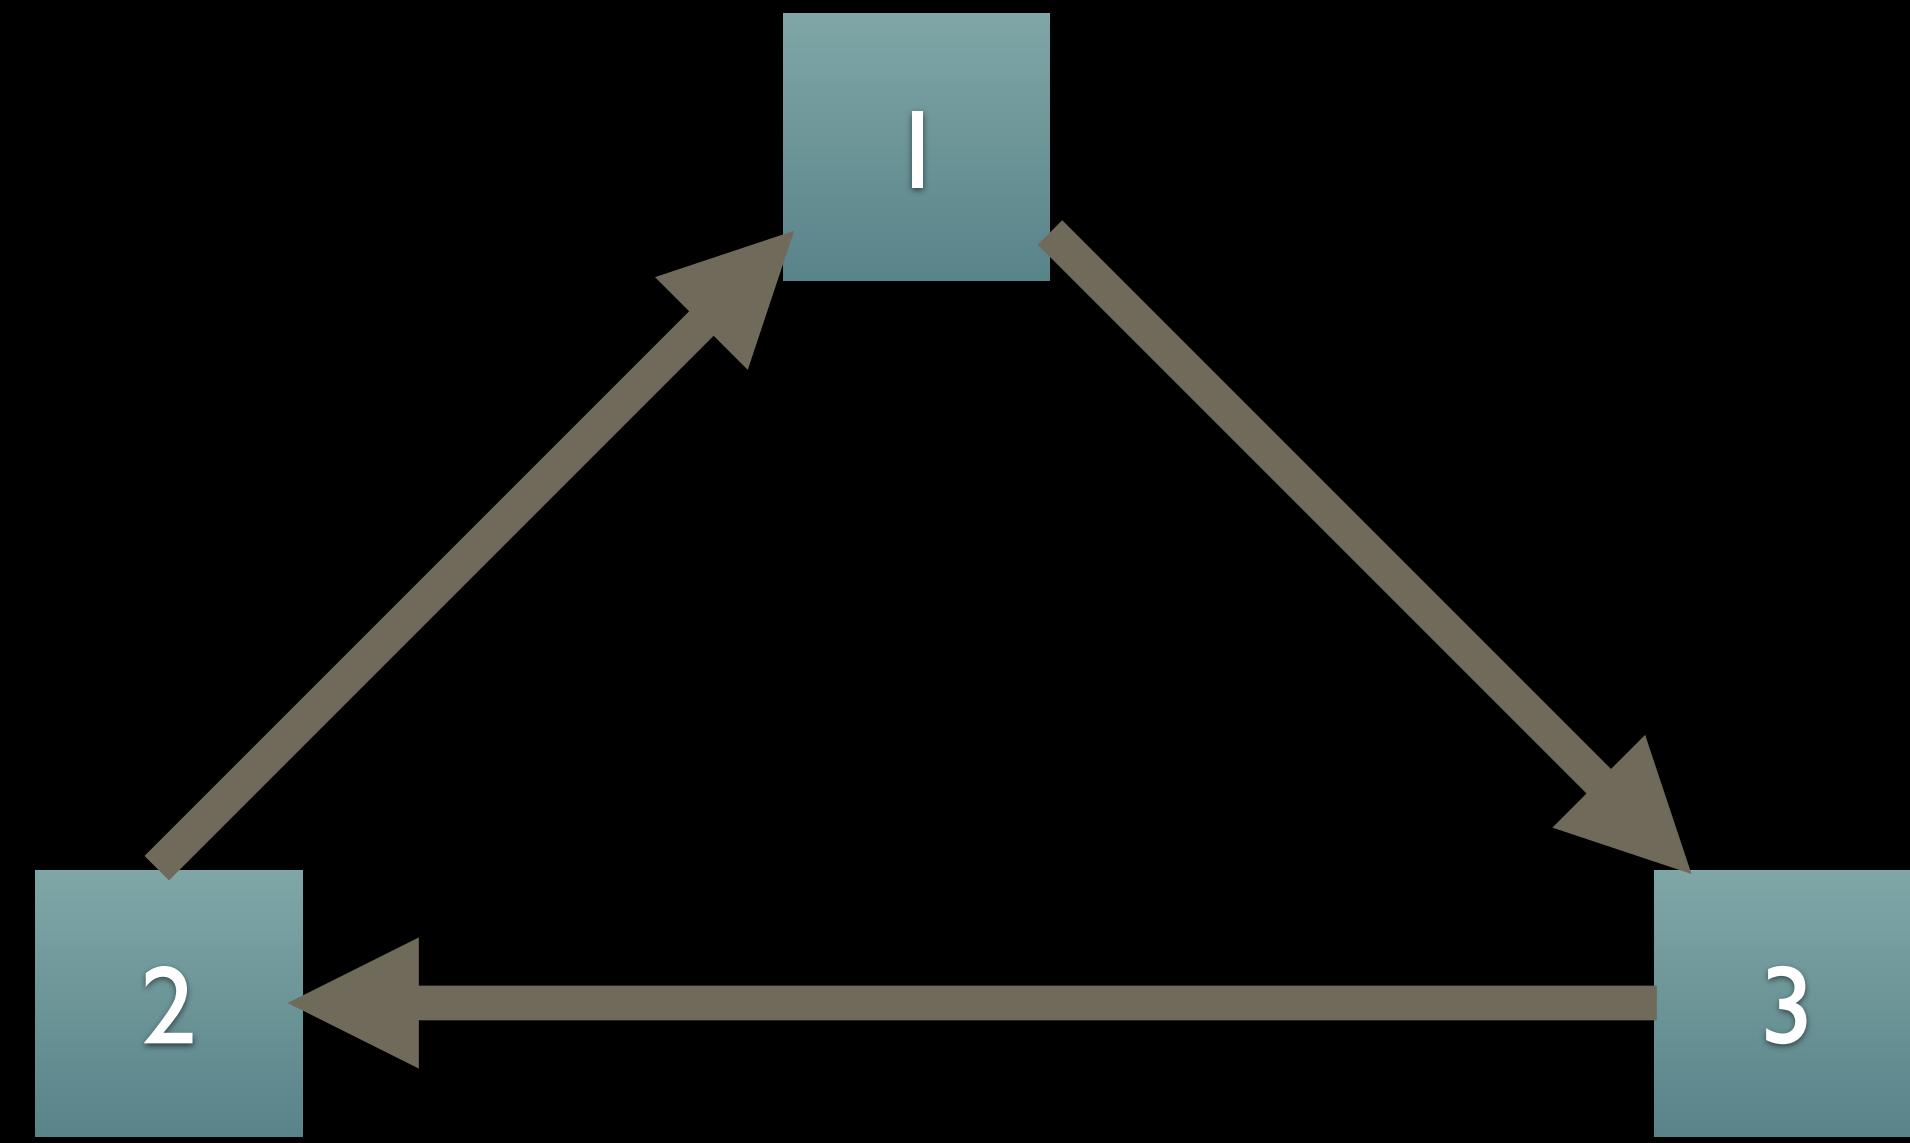
TRANSACTION 3

ASYNC WORKER
Select * from users





LOAD TEST YOUR MIGRATIONS.



scaling with your
team

Migration User

Migration Tooling

Load Testing

Object Relational
Mapping(ORM)

offload to your object relation model (ORM)

- Source of Truth

offload to your object relation model (ORM)

- Source of Truth
- Add sensible defaults to your ORM

offload to your object relation model (ORM)

- Source of Truth
- Add sensible defaults to your ORM
- Every model in every language needs to be constantly updated

remaining issues
getting to zero
downtime

Foreign Keys :(
Vacuuming
Primary and
Replica

foreign keys :(

Alter Table Orders

Add constraint "user_fk"

foreign key("user_id") references "users" ("id")

Both ORDERS and USERS acquires an ACCESS EXCLUSIVE
LOCK!

Both these tables can't be read or written to while this is happening.

DO YOU REALLY NEED REFERENTIAL INTEGRITY?

vacuuming

You can use https://github.com/reorg/pg_repack to vacuum tables without running a VACUUM FULL.

Vacuuming does lock the table to reads and writes and will create downtime otherwise.

primary and replica

As you scale out your primary and replica instances, you want to move all your reads to the replica.

This means fast processing time on the primary = less likely to have long locking schema changes

scaling with tools

Break your
migrations

Don't lock
yourself out

Migration User
Migration Tooling

scaling with your team

remaining
issues
getting to
zero
downtime

Foreign Keys :(

Jumana Bahrainwala
jzbahrai.github.io
@JumzB





Jumana Bahrainwala
jzbahrai.github.io
[@JumzB](https://twitter.com/JumzB)

Can't sing but will rap the
Alexandra Hamilton
song with you



acknowledgments

- Gary Brenhardt, Brent Sullivan, Nicolas Neu, Filipe Fernandes, Alex Snurnikov, Zach Howard, Dileshni Jayasinghe and all the other wonderful folk @Unata
- Braintree - Safe operations for High volume postgresql
- Nordeus Engineering - PSQL Locking
- Travis of the North - PSQL adding FK with zero downtime
- Lob - Running Vacum full with minimum downtime
- Sam Saffron - Managing DB schema Changes without downtime at Discord
- Citus - When PSQL blocks
- Postgresql - Explicit Locking

WHAT'S THAT?

NO IDEA. WIKIPEDIA IS
DOWN, SO I CAN'T CHECK.



