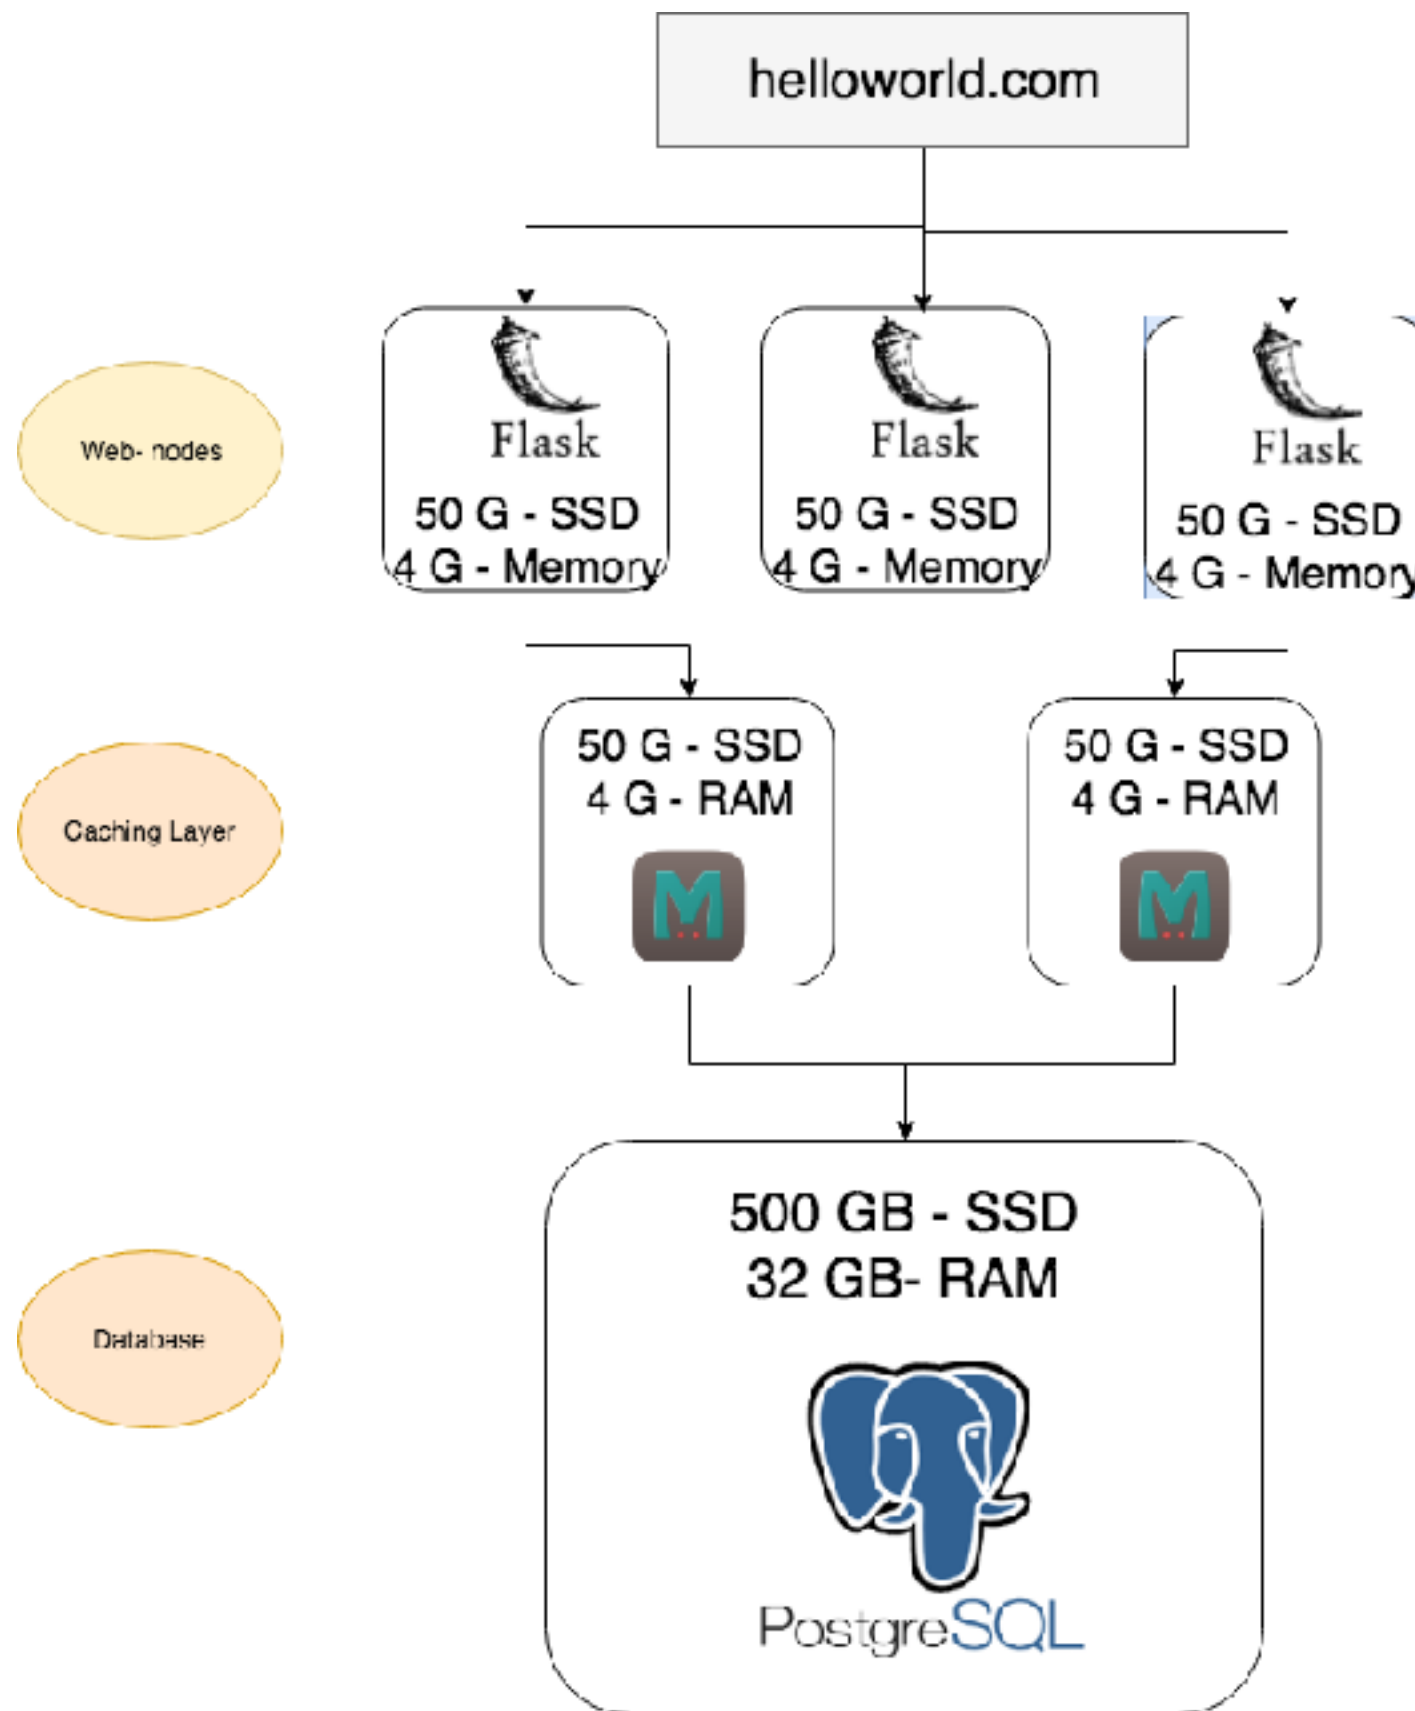


---


# HOW TO SCALE DATABASE MIGRATIONS

---


Jumana Bahrainwala, April 2018  
@JumzB



**WHEN YOU CHANGE A COLUMN TYPE**



**AND LOCK YOUR DB FOR 10 MINUTES**



---

scaling with tools

scaling with people

remaining issues getting to zero  
downtime

---

---

# scaling with tools

Don't lock yourself  
out

Break your  
migrations

Database + code  
compatibility

---

---

# don't lock yourself out

---

Locks are the core of most Relational Databases.

Locks help with the concurrency and keeping data consistent in an RDBMS.

Locks are great... till they are NOT...especially when running migrations

---

---

“Look for locks that are conflicting.  
Conflicting locks are what cause downtime”

---



---

# don't lock yourself out

---

ACCESS SHARE - Select

---

# don't lock yourself out

---

**ACCESS SHARE** - Select

**ROW SHARE** - Select for Update, Select for Share

---

# don't lock yourself out

---

**ACCESS SHARE** - Select

**ROW SHARE** - Select for Update, Select for Share

**ROW EXCLUSIVE** - Update, Delete, Insert. Anything that modifies Data.

---

---

# don't lock yourself out

---

**ACCESS SHARE** - Select

**ROW SHARE** - Select for Update, Select for Share

**ROW EXCLUSIVE** - Update, Delete, Insert. Anything that modifies Data.

**ACCESS EXCLUSIVE** - Alter Table, Drop Table, Truncate, Reindex, Cluster, Vacuum

---

---

# break your migrations

---

“Schema - organization of data as a blueprint of how the database is constructed”

---

# break your migrations

---

“Schema - organization of data as a blueprint of how the database is constructed”

## Schema Changes

- Create, Alter commands change the underlying Data Definition of the DB. They lock the table completely
  - Want these migrations to be as quick as possible and in small transactions
  - **ACCESS EXCLUSIVE**
-

---

# break your migrations

---

“An **index** is a copy of selected columns of **data** from a table that can be searched very efficiently that also includes a low-level disk block address or direct link to the complete row of **data** it was copied from.” - Wikipedia

---

# break your migrations

---

“An **index** is a copy of selected columns of **data** from a table that can be searched very efficiently that also includes a low-level disk block address or direct link to the complete row of **data** it was copied from.” - Wikipedia

## Index Changes

- Adding an index can be done concurrently without locking a table (Postgres)
  - If you are adding an index non-concurrently - recognize it will lock the table
  - **SHARE LOCK**
-



---

# break your migrations

---

Reading from a table and writing to another can be long, especially during a migration.

---

---

# break your migrations

---

Reading from a table and writing to another can be long, especially during a migration.

## Data Migrations

- Chunk your data migrations
  - And insert, update statements should all be in their own transactions as they are non-locking but long.
  - ROW EXCLUSIVE
-

---

Back to Locking in theory

---

---

# don't lock yourself out

---

Alter Table ORDERS

Alter column "final\_order\_total"  
type Numeric(10, 4)

Acquires an ACCESS EXCLUSIVE LOCK

---

---

# don't lock yourself out

Select \* from Orders

ACCESS SHARE

Insert into Orders(amount)

Values(2)

ROW EXCLUSIVE

Both these locks conflict with ACCESS EXCLUSIVE

If the alter statement takes a long time as your orders table is large, your customers **will not be able to READ or WRITE** to the table at that point.

---

---

# don't lock yourself out: ways to get around locking

---

Changing the type of a column

Add a new column with new type (Schema)

Write to both columns and then backfill new column. (Data)

New version of the code stops writing to the old column.  
(Process)

---

---

# don't lock yourself out: ways to get around locking

---

Add a column that has a default:

Add column (SCHEMA)

Add default as a separate command (DATA)

Backfill the column with the default value (DATA)

---

---

# don't lock yourself

---

Add an Index

Add the index using the  
“CONCURRENTLY” keyword.

Add a column that is non-nullable

Create a new table with the addition of  
the non-nullable column, write to both  
tables, backfill, and then switch to the  
new table

Add a column with a unique  
constraint

Add Column, Add unique index  
concurrently and then add the constraint  
on the table.

---



---

# database + code compatibility

---

- Your code should be backwards AND forward compatible
  - Read from BOTH old columns and new
- Don't drop columns until multiple release versions ahead

---

# scaling with your team

Migration User

Migration Tooling

Load Testing

Object Relational  
Mapping(ORM)

---

# migration - user account

---

Make a user to run your migrations

If migrations can't acquire a lock in a timely manner, other statements won't be stuck behind it.

```
ALTER ROLE jzbahrai  
SET lock_timeout='5s';
```

---

---

USERS TABLE

---

---

SELECT

SELECT

USERS TABLE

---

---

ALTER

SELECT

SELECT

USERS TABLE

---

---

# migration - user account

---

SELECT

SELECT

ALTER

SELECT

SELECT

USERS TABLE

---

---

SELECT

SELECT

---

SELECT

SELECT

SELECT

SELECT

SELECT

SELECT

SELECT

USERS TABLE

---



---

# migration tooling

---

1) Automate Scripts:  
No one should be writing Create statements by going to your DB.

# migration tooling

version\_num

-----

39b980e419a9

7ac4c29288a3

16af97dd95c1

24187609b002

9d3ef281fb30

8918b762b17b

47f9fc04d6a7

2badb70854a9

1551e5d79984

## 1) Automate Scripts

No one should be writing Create statements by going to your DB

2) Have a way to know exactly what state your DB is in. Keep a **hash table in your DB** to indicate which migrations have run.

# migration tooling

---

version\_num

-----

39b980e419a9

7ac4c29288a3

16af97dd95c1

24187609b002

9d3ef281fb30

8918b762b17b

47f9fc04d6a7

2badb70854a9

1551e5d79984

## 1) Automate Scripts

No one should be writing Create statements by going to your DB

2) Have a way to know exactly what state your DB is in. Keep a **hash table in your DB** to indicate which migrations have run.

3) Easy set up on dev machines

# migration tooling

version\_num

-----

39b980e419a9

7ac4c29288a3

16af97dd95c1

24187609b002

9d3ef281fb30

8918b762b17b

47f9fc04d6a7

2badb70854a9

1551e5d79984

## 1) Automate Scripts

No one should be writing Create statements by going to your DB

2) Have a way to know exactly what state your DB is in. Keep a **hash table in your DB** to indicate which migrations have run.

3) Easy set up on dev machines

4) Use existing tooling

---

# load test your migrations (demo)

---

- Always run migrations during low traffic

---

# load test your migrations (demo)

---

- Always run migrations during low traffic
- What if low traffic still means you have 1000's of users on your site?

---

# load test your migrations (demo)

---

- Always run migrations during low traffic
  - What if low traffic still means you have 1000's of users on your site?
    - Load test your migration in stage to see where your application gets affected.
    - Go back and either break up your migration further or ask for a maintenance window
-

---

# offload to your object relation model (ORM)

---

- Is your database the source of truth of your data or the ORM?



---

# offload to your object relation model (ORM)

---

- Is your database the source of truth of your data or the ORM?
- You can add sensible defaults to your ORM without having to let anyone go through SQL

---

# offload to your object relation model (ORM)

---

- Is your database the source of truth of your data or the ORM?
  - You can add sensible defaults to your ORM without having to let anyone go through SQL
  - If you do use your ORM as the source of truth:
    - Every model in every language needs to be constantly updated
    - Your ORM needs to be backwardly compatible with your code
-

---

remaining issues  
getting to zero  
downtime

Foreign Keys :(

Vacuuming

Slave and Master

---

---

# foreign keys :(

---

```
Alter Table Orders  
Add constraint "user_fk"  
foreign key("user_id") references  
"users" ("id")
```

Both ORDERS and USERS acquires an ACCESS EXCLUSIVE LOCK!

Both these tables can't be read or written too while this is happening.

DO YOU REALLY NEED REFERENTIAL INTEGRITY?

---

---

# vacuuming

---

You can use [https://github.com/reorg/pg\\_repack](https://github.com/reorg/pg_repack) to vacuum tables without running a VACCUUM FULL.

Vacuuming does lock the table to reads and writes and will create downtime otherwise.

---

---

# slave and master

---

As you scale out your slave and master instances, you want to move all your reads to the slave.

This means fast processing time on the master = less likely to have long locking schema changes

---

---

scaling with  
tools

remaining  
issues  
getting to  
zero  
downtime

Break your  
migrations

Migration User  
Migration Tooling

Don't lock  
yourself out

Foreign  
Keys :(

scaling with your  
team

---

---

# acknowledgments

---

- Nicolas Neu, Filipe Fernandes, Alex Snurnikov, Zach Howard, Dileshni Jayasinghe and all the other wonderful folk @Unata
  - Nordeus Engineering - PSQL Locking
  - Travis of the North - PSQL adding FK with zero downtime
  - Lob - Running Vaccum full with minimum downtime
  - Sam Saffron - Managing DB schema Changes without downtime at Discord
  - Citus - When PSQL blocks
  - Postegresql - Explicit Locking
-



WHAT'S THAT?

NO IDEA. WIKIPEDIA IS  
DOWN, SO I CAN'T CHECK.



Requested Lock Mode	Current Lock Mode							
	ACCESS SHARE	ROW SHARE	ROW EXCLUSIVE	SHARE UPDATE EXCLUSIVE	SHARE	SHARE ROW EXCLUSIVE	EXCLUSIVE	ACCESS EXCLUSIVE
ACCESS SHARE								X
ROW SHARE							X	X
ROW EXCLUSIVE						X	X	X
SHARE UPDATE EXCLUSIVE					X	X	X	X
SHARE				X	X	X	X	X
SHARE ROW EXCLUSIVE			X	X	X	X	X	X
EXCLUSIVE		X	X	X	X	X	X	X
ACCESS EXCLUSIVE	X	X	X	X	X	X	X	X